COMPUTER SCIENCE

Database Management System

Transaction & Concurrency Control

Lecture_2

Vijay Agarwal sir

01 Serializable Schedule

02 Conflict & View Serializable

Transaction Concept

Transaction State

ACID

A: Atomicity

C: Consistency

I: Isolation

D: Durability

Schedule

- Serial Schedule (n!)
- Non Serial Schedule

Note: ALL Serial Schedule are always Consistent.

Note: Non serial schedule May (or) May Not be Consistent.

But we execute Non Serial Schedule (Concurrent Execution).

Let T_1 transfer 100 Rs from A to B, and T_2 transfer 10% of the balance from A to B.

A: 2000
+B: 3000
5000

Schedule 1

T_1	T_2
read (A)	
A: = A – 100	
write (A)	
read (B)	
B: = B + 100	
write (B)	
commit	
	read (A)
	temp := A * 0.1
	A := A – temp
	write (A)
	read (B)
	B := B + temp
	write (B)
	Commit

A: 1710
+B: 3290
5000
Consistent

$S_1 < T_1\ T_2 >$

Serial schedule in which T_1 is followed by T_2:

Schedule 2

T_1	T_2
	read (A)
	temp := A * 0.1
	A := A – temp
	write (A)
	read (B)
	B := B + temp
	write (B)
	Commit
read (A)	
A: = A – 100	
write (A)	
read (B)	
B: = B + 100	
write (B)	
commit	

$S_2 < T_2\ T_1 >$

serial schedule where T_2 is followed by T_1

A = 1700
+ B = 3300
5000
Consistent

All Serial Schedule

$S_1 <T_1\ T_2>$ T_1 followed by T_2.

$S_2 <T_2, T_1>$ T_2 followed by T_1

are always Consistent.

Schedule 3

T_1	T_2
read (A)	
A: = A – 100	
write (A)	
	read (A)
	temp := A * 0.1
	A := A – temp
	write (A)
read (B)	
B: = B + 100	
write (B)	
commit	
	read (B)
	B := B + temp
	write (B)
	Commit

A: 1710
+ B: 3290
5000

Consistent

$C_1 \equiv S_1 (T_1, T_2)$

Schedule 4

T_1	T_2
read (A)	
A: = A – 100	
	read (A)
	temp := A * 0.1
	A := A – temp
	write (A)
write (A)	
read (B)	
B: = B + 100	
write (B)	
commit	
	read (B)
	B := B + temp
	write (B)
	Commit

A: 1900
+ B: 3300
5200

Inconsistent

C_2

(C1) ✓ ≡ S1 ⟨T1, T2⟩

Non Serial Schedule C_1 is Consistent & its equivalent to serial schedule S_1 ⟨T1, T2⟩ T1 followed by T2.

Non Serial Schedule C_2 is Not Consistent (or) In Consistent.

C_2 X

Serial Schedule

- After Commit of one transaction, begins (Start) another transaction.
- Number of possible serial Schedules with 'n' transactions is "n!"
- The execution sequence of Serial Schedule always generates consistent result.

Example

S : $R_1(A)$ $W_1(A)$ Commit (T_1) $R_2(A)W_2(A)$ commit (T_2).

T1 followed by T2

Advantage

- Serial Schedule always produce correct result (integrity guaranteed)

as no resource sharing.

Disadvantage

- Less degree of concurrency.
- Through put of system is low.
- It allows transactions to execute one after another.

serializable schedule

If a Non Serial Schedule (Concurrent Execution) has been executed, that could have Same effect on the Database, as a Schedule executed without Any Concurrent-execution (Any Serial Schedule) is called Serializable Schedule.

(Note) Serializable Schedule are always Consistent.

This Process is Called Serializability.

How to achieve serializable schedule

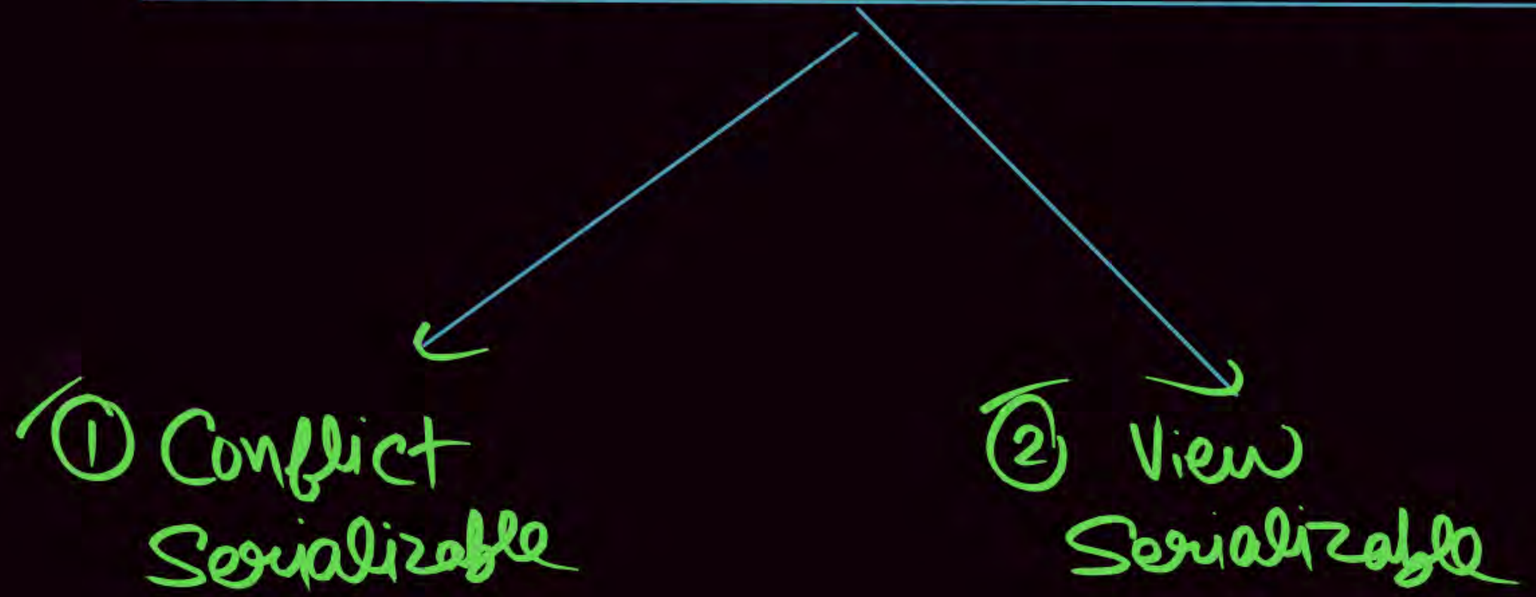

Serializable Schedule

A Schedule is serializable Schedule if it is equivalent to a Serial Schedule.

(i) Conflict Serializability

(ii) View Serializability

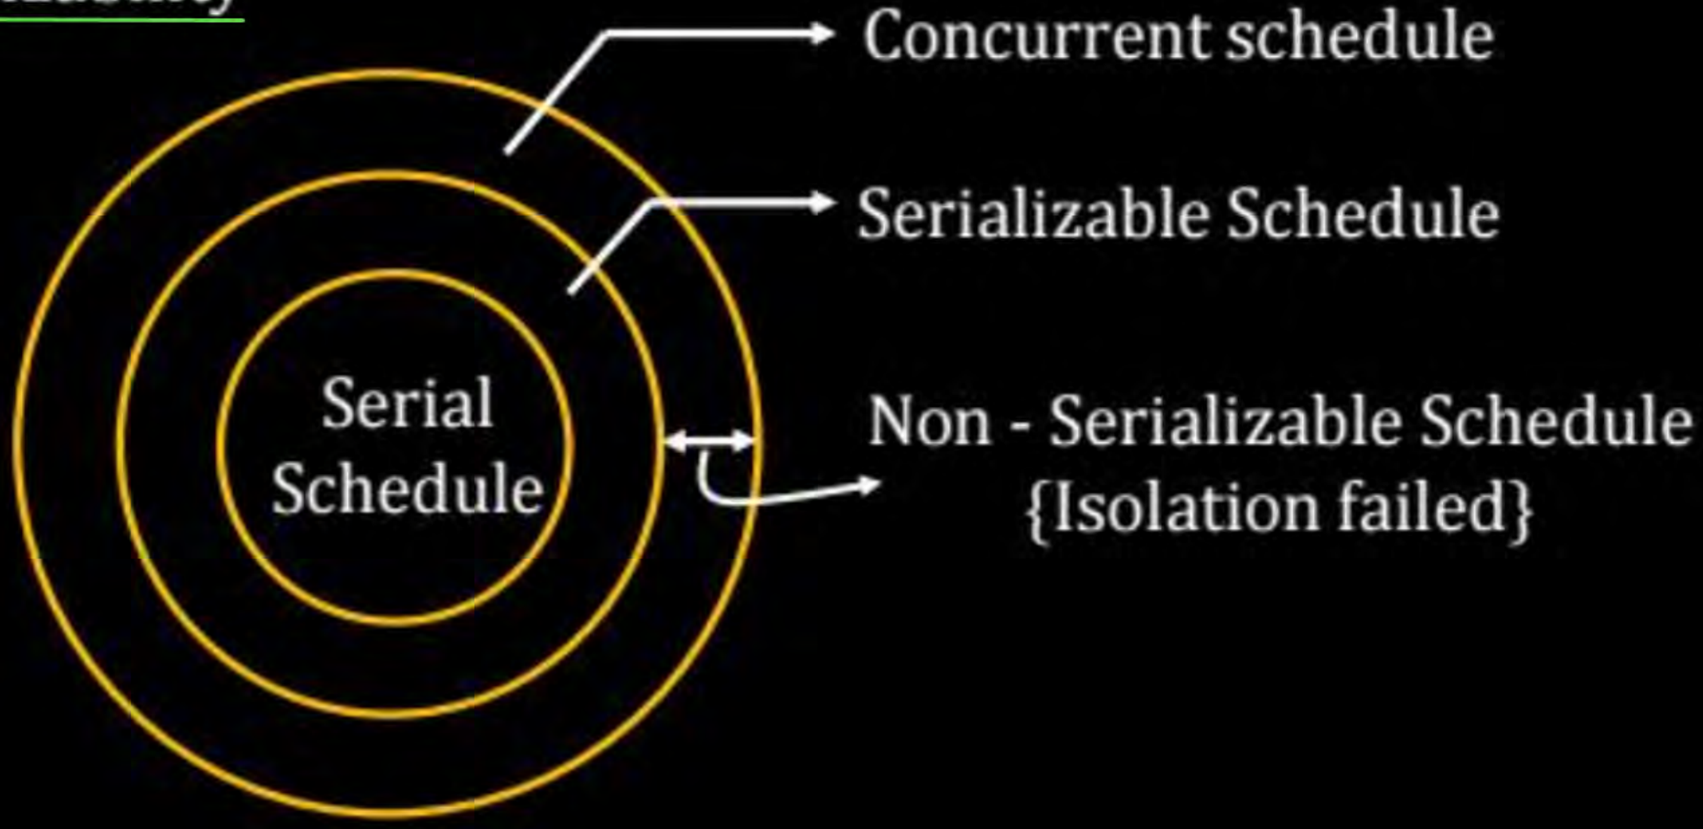

Serializability

- Basic Assumption: Each transaction preserves database consistency.
- Thus, serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
  1. Conflict serializability
  2. view serializability

conflict serializable

1. Basic Concept
2. Testing Method (Precdence Graph) — V.Imp.
3. Conflict Equal to Any Serial Schedule.

conflict serializable

Conflict operation

Let us Consider Schedule S, in which there are Two Consecutive Instruction T_i & I_j of transaction T_i & T_j Respectively $[i \neq j]$

Same Data Item

T_i	T_j
R(A)	W(A)
W(A)	R(A)
W(A)	W(A)

Conflict Instruction/operation [swapping Not possible]

Non Conflict Operation/Instruction

T_i	T_j
R(A)	R(A)
R(A)	W(B)

Different Data Item

Non Conflict Instn/operation [swapping possible]

conflict serializable

S: Non Serial Schedule [Given Question]

S ⟶ (by series of Swap of Non Conflicting Instruction/operation) ⟶ S'

S': Any serial schedule of S' (of Given Question)

then Schedule S is Conflict Serializable.

Conflict Serializability

- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict equivalent.
- We say that a schedule S is conflict serializable if it is conflict equivalent to a serial schedule.

Q

T_1	T_2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

1710
3290
5000

S

C_1

(T_1, T_2)

T_1	T_2
R(A)	
W(A)	
R(B)	
W(B)	
	R(A)
	W(A)
R(B)	
W(B)	
CT	T_2

(T_1, T_2)

Q.1

T_1	T_2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

S

C_1

Not Possible to Convert $S(T_2, T_1)$

T_1	T_2
	R(A)
	W(A)
	R(B)
	W(B)

$S' < T_2, T_1 >$. T_2 followed by T_1.

Q.1

T_1	T_2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)

S

C_1

yes possible to swap & Convert into $S' < T_1\ T_2 >$.

T_1	T_2
R(A)	
W(A)	
R(B)	
W(B)	

T_1	T_2
R(A)	
W(A)	
R(B)	
W(B)	

$S' < T_1, T_2 >$. T_1 followed by T_2.

Ⓖ ing eg.

Here S is Conflict Serializable

to $S' \langle T_1, T_2 \rangle$ [T_1 followed by T_2].

Q.2

T_1	T_2
R(A)	
	R(A)
	W(A)
W(A)	
R(B)	
W(B)	
	R(B)
	W(B)

C2 5200 C2 Inconsistent

Not Possible

T_1	T_2
	R(A)
	W(A)
	R(B)
	W(B)
R(A)	
W(A)	
R(B)	
W(B)	

$< T_2\ T_1 >$

Q.2

T_1	T_2
R(A)	
	R(A)
	W(A)
W(A)	
R(B)	
W(B)	
	R(B)
	W(B)
	C_2

Not possible

T_1	T_2
R(A)	
W(A)	
R(B)	
W(B)	
	R(A)
	W(A)
	R(B)
	W(B)

$S'(T_1, T_2)$

C_2 is Not Possible to convert either

$S' < T_1, T_2 >$

& $S' < T_2, T_1 >$

So C_2 is Not Conflict Serializable.

C_2

Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1, by series of swaps of non-conflicting instructions. Therefore Schedule 3 is conflict serializable.

Schedule 3

T_1	T_2
read (A)	
Write (A)	
	read (A)
	write (A)
read (B)	
write (B)	
	read (B)
	write (B)

Schedule 6

T_1	T_2
read (A)	
write (A)	
read (B)	
write (B)	
	read (A)
	write (A)
	read (B)
	write (B)

Conflict Serializability (Cont.)

- Example of a schedule that is not conflict serializable:

T_3	T_4
read (Q)	
	write (Q)
write (Q)	

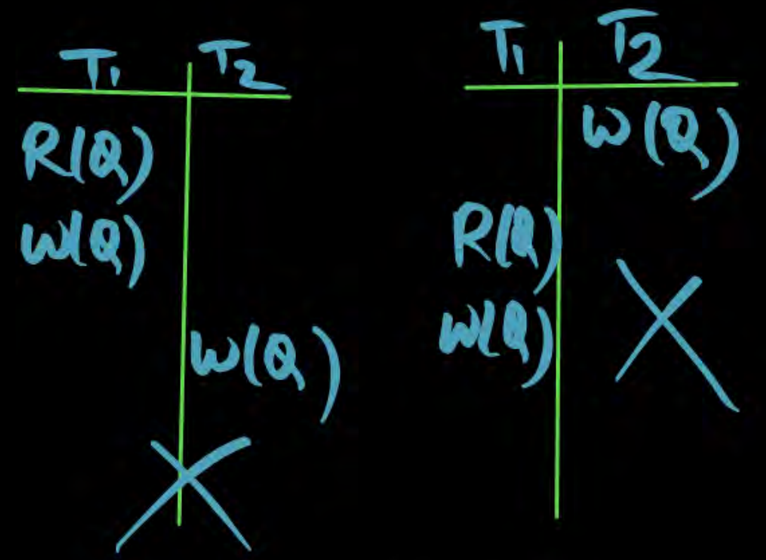

- We are unable to swap instructions in the above schedule to obtain either the serial schedule < T3, T4 >, or the serial schedule < T4, T3 >

Conflict Serializable

A schedule is said to be conflict serializable if it is conflict equivalent to a serial schedule.

Same conflicting operation order in C_1 & S_1

$\therefore$ Its $\{C_1\}$ conflict is conflict serializable.

T_1	T_2
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

C_L

T_1	T_2
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)

S_L

Conflicting Instructions

- Instructions I_i, and I_j of transactions T_i and T_j respectively, conflict if and only if there exists some item Q accessed by both I_i, and I_j, and at least one of these instructions wrote Q. ($i \neq j$)
  1. I_i, = read(Q), I_j = read(Q). I_i and I_j don't conflict. (Non Conflict Instⁿ)
  2. I_i, = read(Q) I_j = write(Q). They conflict.
  3. I_i, = write(Q) I_j = read(Q). They conflict
  4. I_i= write(Q) I_j = write(Q). They conflict

  (2, 3, 4: Conflict Instruction)
- Intuitively, a conflict between I_i and I_j forces a (logical) temporal order between them.
  - If I_i, and I_j are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

Testing for conflict serializable

Precedence Graph Method

$G(V, E)$

(vertex) V: Set of Transactions.

(Edge) E: Edge $T_i \longrightarrow T_j$ edge occur iff any one Condition Hold

T_i T_j

Conflict operation
- $R(A) \longrightarrow W(A)$
- $W(A) \longrightarrow R(A)$
- $W(A) \longrightarrow W(A)$

then edge Draw.

Testing for conflict serializable

Precedence Graph Method

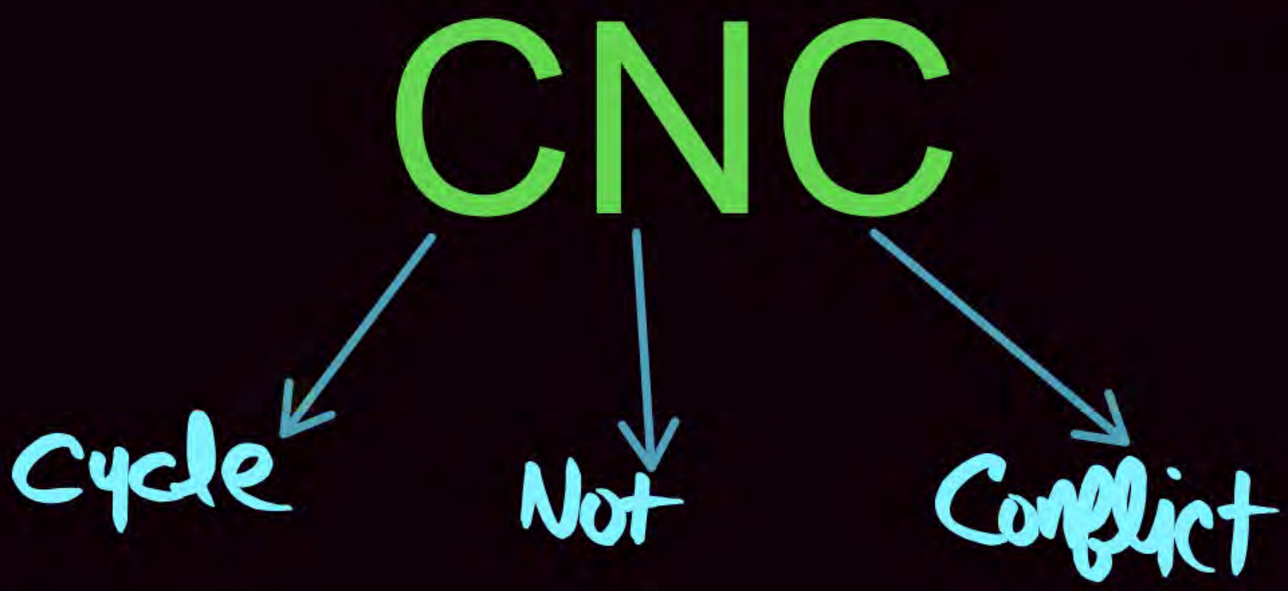

(Note) If Precedence Graph Contain cycle [Any one cycle] then schedule is Not Conflict serializable

Testing for Serializability

- Testing for conflict serializability.
  - Consider some schedule of a set of transactions T_1, T_2, ...T_n
  - Precedence graph — a direct graph where the vertices are the transactions (names).
  - We draw an arc from T_i to T_j if the two transaction conflict, and T_i accessed the data item on which the conflict arose earlier.
  - We may label the arc by the item that was accessed.

Example:

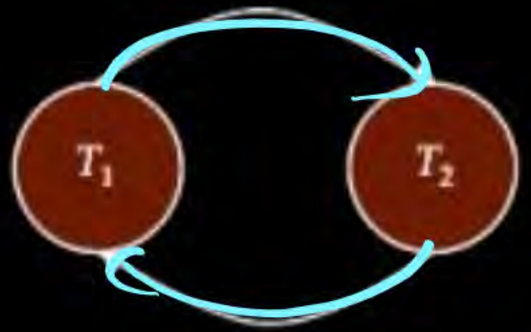

CNC

Cycle Not Conflict Serializable.

A schedule is conflict serializable if and only if its precedence graph is acyclic.

Conflict serializable

NOTE: CNC [Cycle not conflict serializable]

Q.1

S: $R_1(A)\ W_1(A)\ R_2(A)\ W_2(A)\ R_1(B)\ W_1(B)\ R_2(B)\ W_2(B)$

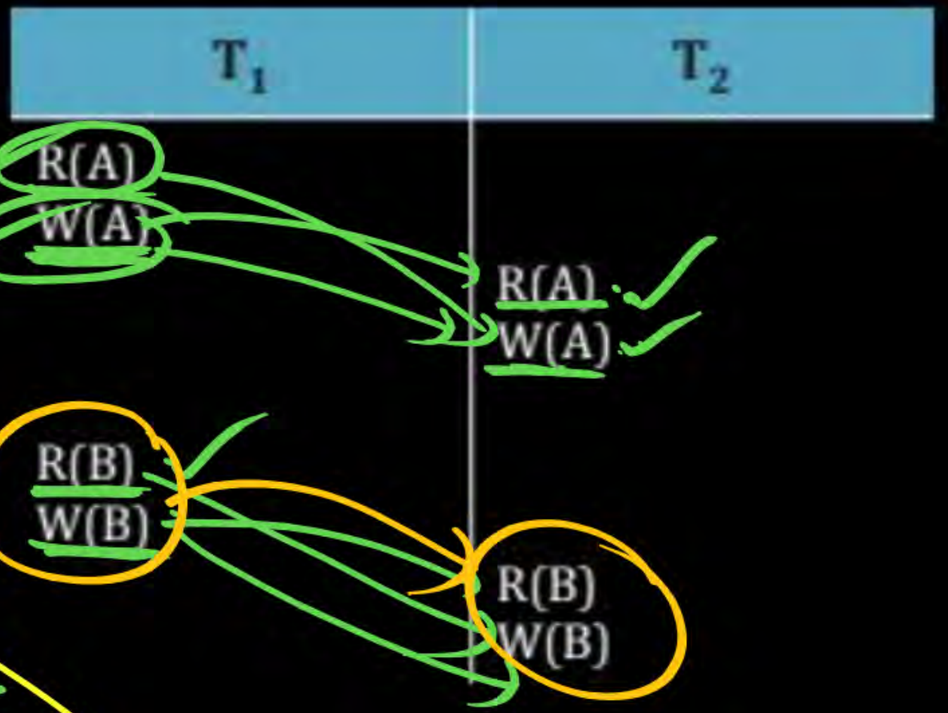

Acyclic

Conflict Serializable

$\langle T_1, T_2 \rangle$

C1

1710
3290
5000

Consistent

$\langle T_1, T_2 \rangle$

Q.2 $R_1(A)\ R_2(A)\ W_2(A)\ W_1(A)\ R_1(B)\ W_1(B)\ R_2(B)\ W_2(B)$

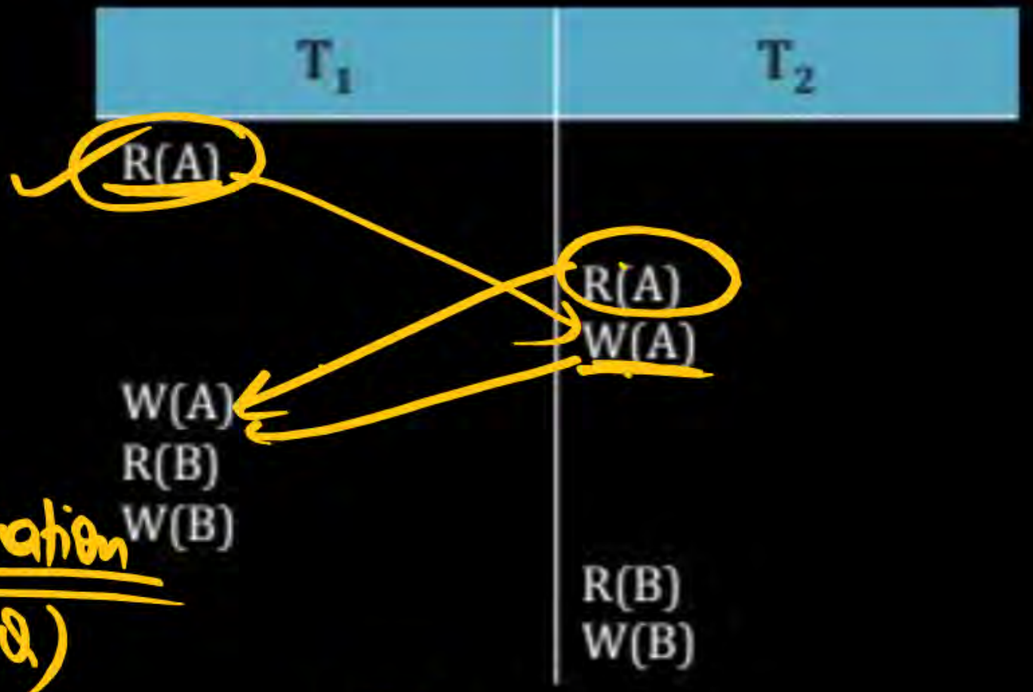

Conflict operation

R(Q) → W(Q)

W(Q) → R(Q)

W(Q) → W(Q)

Cycle Not Conflict

5200 Inconsistent

Q.3

$R_1(A)\ W_1(A)\ R_2(B)\ W_2(B)\ R_1(B)\ W_1(B)\ R_2(A)\ W_2(A)$

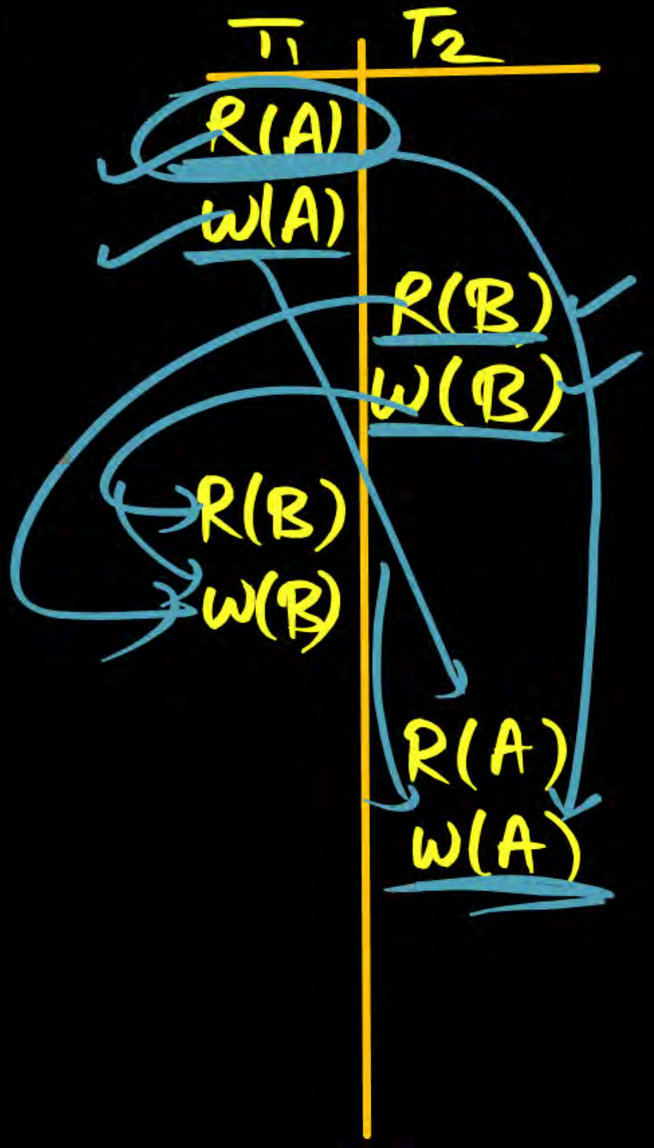

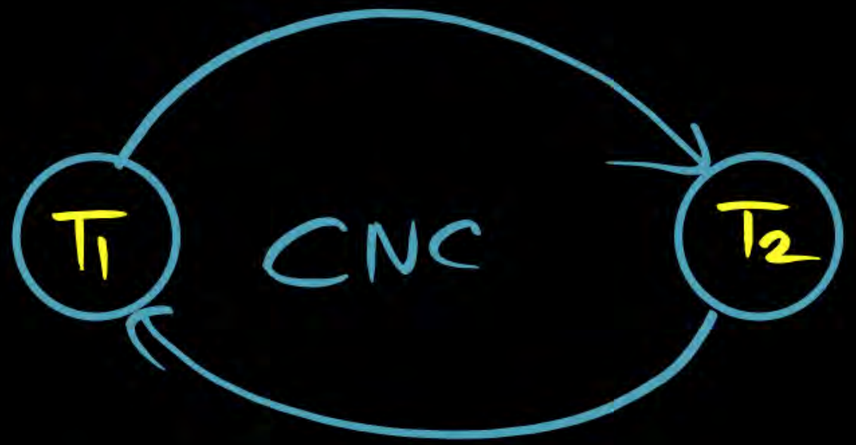

Cycle Not Conflict

Q.1

Consider the following schedules involving two transactions. Which one of the following statements is TRUE?

S_1: $r_1(X); r_1(Y); r_2(X); r_2(Y); w_2(Y); w_1(X)$

S_2: $r_1(X); r_2(X); r_2(Y); W_2(Y); r_1(Y); w_1(X)$

[2007: 2 Marks]

A. Both S_1 and S_2 are conflict serializable

B. S_1 is conflict serializable and S_2 is not conflict serializable

C. S_1 is not conflict serializable and S_2 is conflict serializable

D. Both S_1 and S_2 are not conflict serializable

S_1: $r_1(x)\ r_1(y)\ r_2(x)\ r_2(y)\ w_2(y)\ w_1(x)$

S_2 $r_1(x)\ r_2(x)\ r_2(y)\ w_2(y)\ r_1(y)\ w_1(x)$

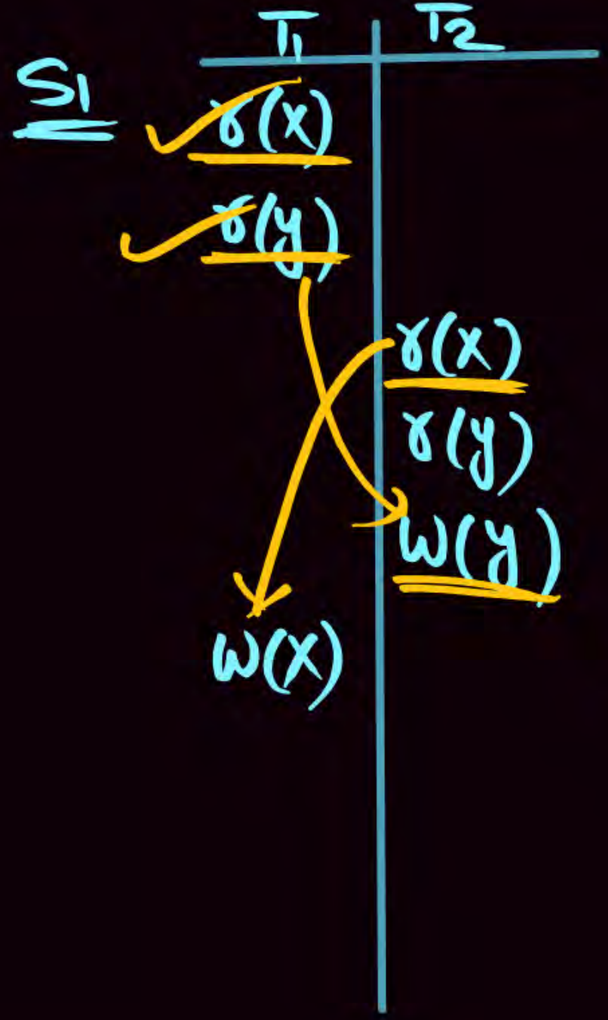

Cycle Not Conflict

S_1 is Not Conflict

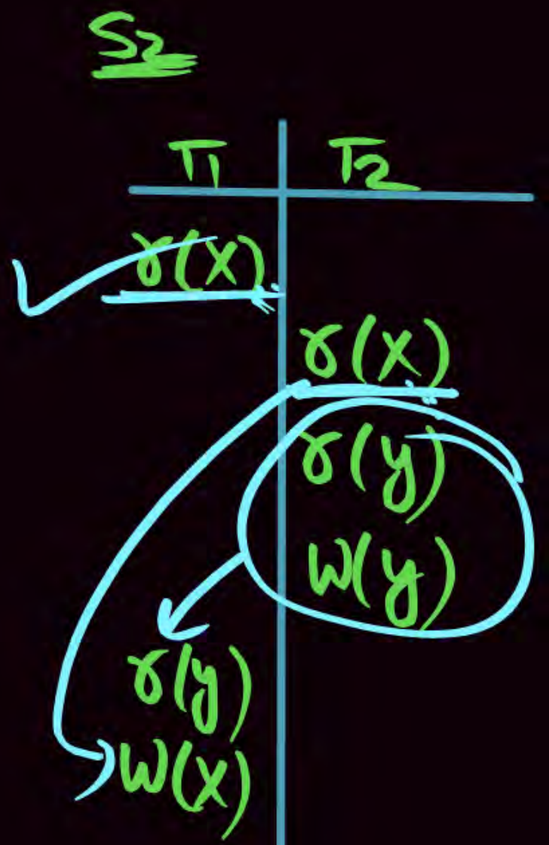

$\langle T_2, T_1 \rangle$

S_2 Conflict Serializable

Q.2

Consider the following four schedules due to three transactions (indicated by the subscript) using read and write on a data item x, denoted by r(x) and w(x) respectively. Which one of them is conflict serializable?

[2014(Set-1): 2 Marks]

A $r_1(x); r_2(x); w_1(x); r_3(x); w_2(x)$

B $r_2(x); r_1(x); w_2(x); r_3(x); w_1(x)$

C $r_3(x); r_2(x); r_1(x); w_2(x); w_1(x)$

D $r_2(x); w_2(x); r_3(x); r_1(x); w_1(x)$

MCQ Q.3

Consider the transactions T1, T2 and T3 and the schedules S1 and S2 given below.

T1: r1(X); r1(Z); w1(X); w1(Z)
T2: r2(Y) ; r2(Z) ; w2(Z)
T3: r3(Y); r3(X); w3(Y)

S1: r1(X); r3(Y); r3(X); r2(Y); r2(Z); w3(Y); w2(Z); r1(Z); w1(X); w1(Z)
S2: r1(X); r3(Y); r2(Y); r3(X); r1(Z); r2(Z); w3(Y); w1(X); w2(Z); w1(Z)

Which one of the following statements about the schedules is TRUE?

[GATE-2014-CS: 2M]

A. Only S1 is conflict-serializable.

B. Only S2 is conflict-serializable.

C. Both S1 and S2 are conflict-serializable.

D. Neither S1 nor S2 is conflict-serializable.

Q.4

Let $r_i(z)$ and $w_i(z)$ denote read and write operations respectively on a data item by a transaction T_i. Consider the following two schedules.

S_1: $r_1(x)$ $r_1(y)$ $r_2(x)$ $r_2(y)$ $w_2(y)$ $w_1(x)$

S_2: $r_1(x)$ $r_2(x)$ $r_2(y)$ $w_2(y)$ $r_1(y)$ $w_1(x)$

Which one of the following options is correct? [MCQ: 2021: 2M]

A. S_1 is conflict serializable, and S_2 is not conflict serializable.

B. S_1 is not conflict serializable, and S_2 is conflict serializable.

C. Both S_1 and S_2 are conflict serializable.

D. Neither S_1 nor S_2 is conflict serializable.

Q.5

Let $R_i(z)$ and $W_i(z)$ denote read and write operations on a data element z by a transaction T_i, respectively. Consider the schedule S with four transactions.

S: $R_4(x), R_2(x), R_3(x), R_1(y), W_1(y), W_2(x), W_3(y), R_4(y)$

Which one of the following serial schedules is conflict equivalent to S?

[2022: 2 Marks]

A $T_1 \rightarrow T_3 \rightarrow T_4 \rightarrow T_2$

B $T_1 \rightarrow T_4 \rightarrow T_3 \rightarrow T_2$

C $T_4 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2$

D $T_3 \rightarrow T_1 \rightarrow T_4 \rightarrow T_2$

Q.6

Consider the following transaction involving two bank accounts x and y.

read(x); x: = x – 50; write (x); read (y); y: = y + 50; write (y)

The constraint that the sum of the accounts x and y should remain constant is that of [2015(Set-2): 1 Marks]

A Atomicity

B Consistency

C Isolation

D Durability

MCQ Q.7

Which one of the following is NOT a part of the ACID properties of database transactions?

[GATE-2016-CS: 1M]

A Atomicity

B Consistency

C Isolation

D Deadlock-freedom

MCQ **Q.8**

Suppose a database schedule S involves transaction T_1....., T_n. Construct the precedence graph of S with vertices representing the transactions and edges representing the conflicts. If S is serializable, which one of the following orderings of the vertices of the precedence graph is guaranteed to yield a serial schedule?

[GATE-2016-CS: 2M]

A Topological order

B Depth-first order

C Breadth-first order

D Ascending order of transaction indices

MCQ Q.9

Consider the following schedule for transactions T1, T2 and T3:

Which one of the schedules below is the correct serialization of the above?

[GATE-2010-CS: 2M]

T1	T3	T3
Read(X)		
	Read (Y)	
		Read (Y)
	Write (Y)	
Write (X)		
		Write (X)
	Read (X)	
	Write (X)	

(A) T 1 → T 3 → T 2

(B) T 2 → T 1 → T 3

(C) T2 → T3 → T1

(D) T3 → T 1 → T 2

MCQ **Q.10**

Consider two transactions T_1 and T_2, and four schedules S_1, S_2, S_3, S_4 of T_1 and T_2 as given below:

T_1: $R_1[x]$ $W_1[x]$ $W_1[y]$;

T_2: $R_2[x]$ $R_2[y]$ $W_2[y]$;

S_1: $R_1[x]$ $R_2[x]$ $R_2[y]$ $W_1[x]$ $W_1[y]$ $W_2[y]$;

S_2: $R_1[x]$ $R_2[x]$ $R_2[y]$ $W_1[x]$ $W_2[y]$ $W_1[y]$;

S_3: $R_1[x]$ $W_1[x]$ $R_2[x]$ $W_1[y]$ $R_2[y]$ $W_2[y]$;

S_4: $R_2[x]$ $R_2[y]$ $R_1[x]$ $W_1[x]$ $W_1[y]$ $W_2[y]$;

Which of the above schedules are conflict serializable?

[GATE-2009-CS: 2M]

A S_1 and S_2

B S_2 and S_3

C S_3 only

D S_4 only

THANK YOU!