

COMPUTER SCIENCE

Database Management System

Transaction & Concurrency Control

Lecture_7



Vijay Agarwal sir

A graphic of a construction barrier with orange and white diagonal stripes and two yellow spherical bollards at the top.

**TOPICS
TO BE
COVERED**

- 01 Lock Based Protocol**

- 02 Time Stamp Protocol**

Serializability Recoverability

Problem due to Concurrent Execution

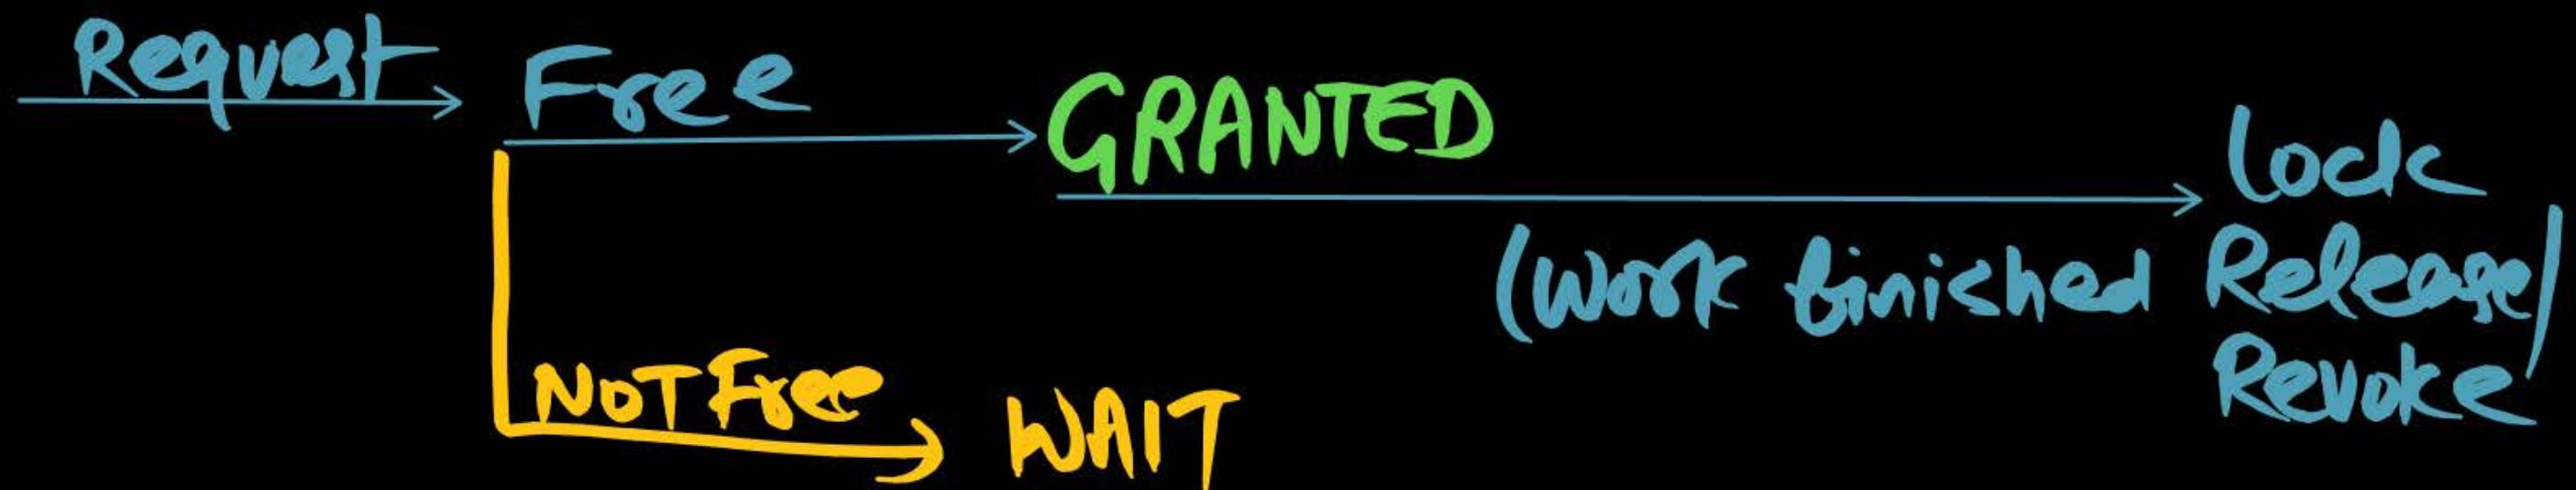
Topological Sorting

Finding Number of Conflict serializable Schedule .

Implementation of Concurrency Control.

LOCK BASED Protocol .

LOCK BASED Protocol



Procedure

Before Using Any Data Item, Transaction should Request for a Lock.
If Lock is Free (Lock Not taken by any other Transaction on that Data Item)
Then it will be granted, otherwise Transaction has to Wait (if Busy)

① SHARED LOCK [S] [Only Read]

② Exclusive lock (X) [write(A) / write(A) / Read(A) / Read(A) / Write(A)]

SHARED Lock [S]

LOCK - S(A)

Read (A)

Unlock - S(A)

Exclusive Lock (X)

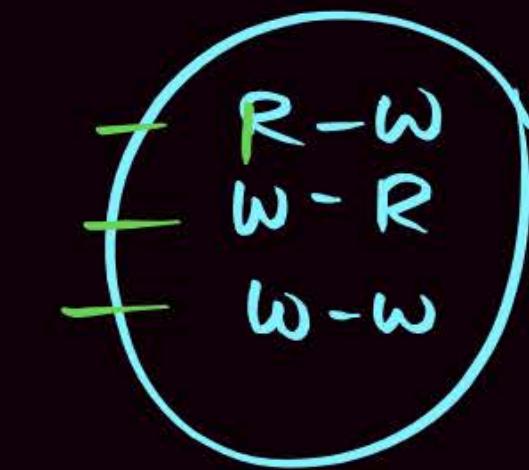
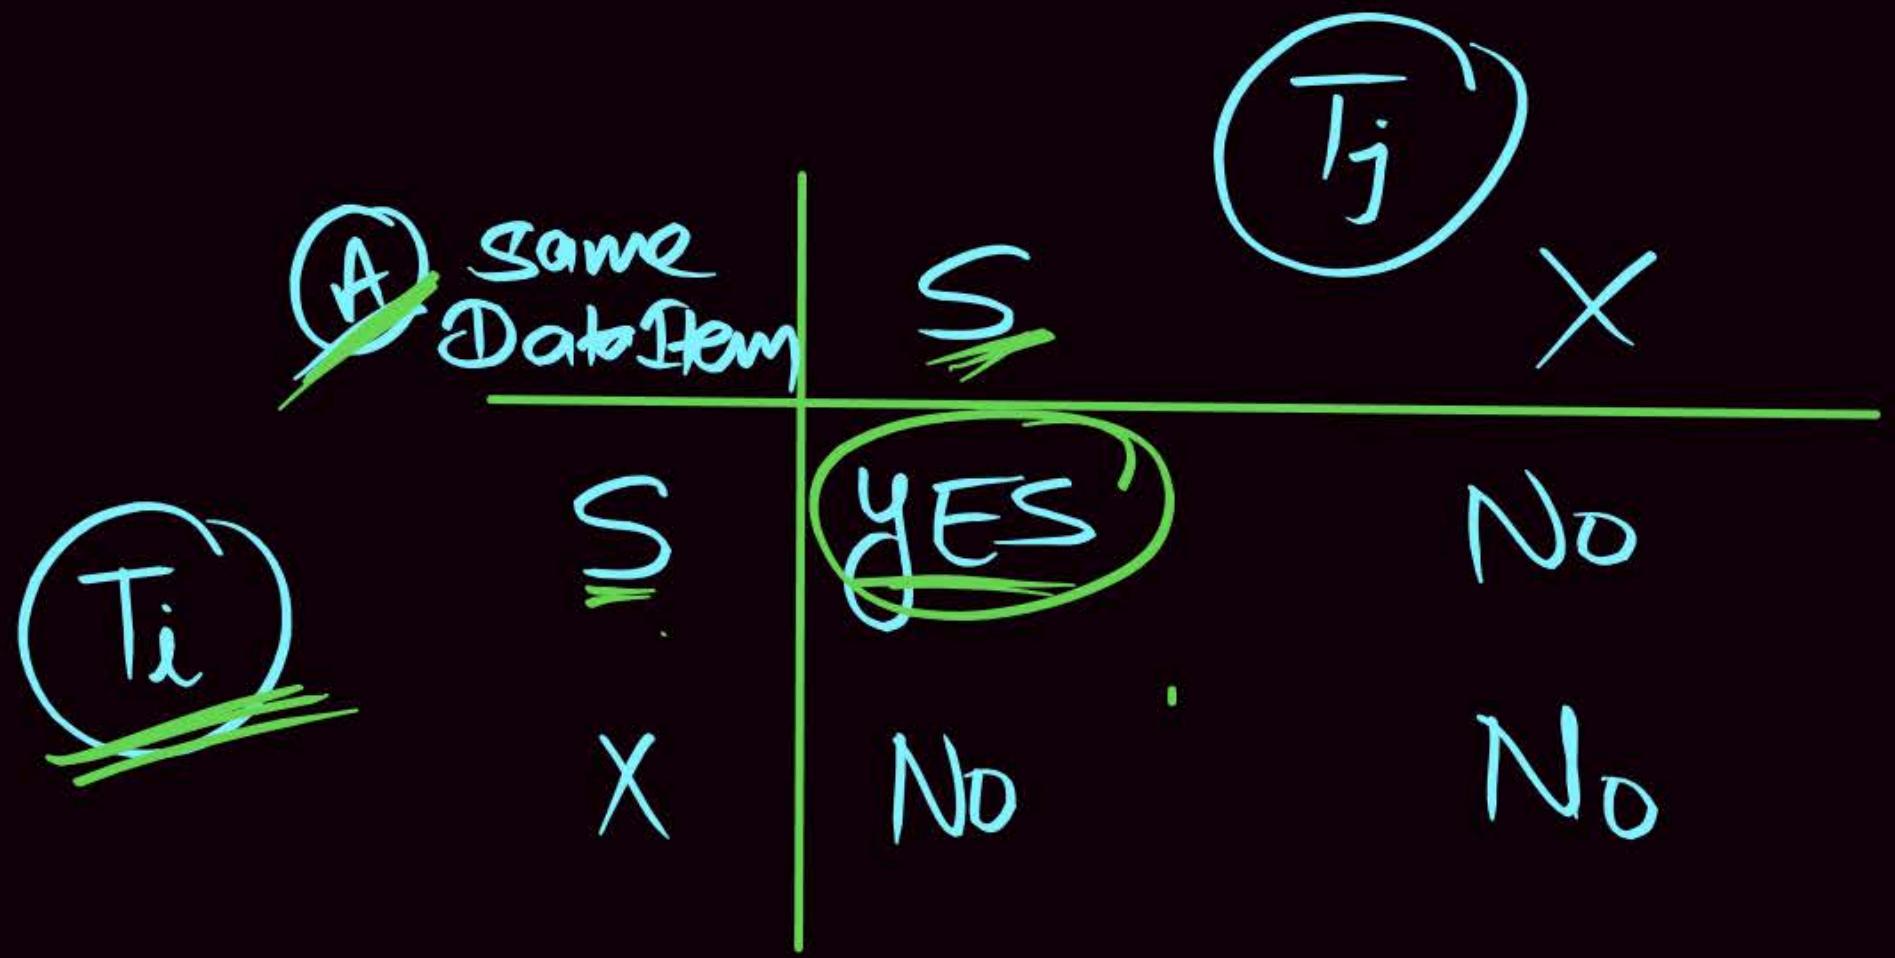
Lock - X(A)

write(A) ⋀ Read(A)
write(A) ⋀ Read(A)

Unlock - X(A)

Lock-Based Protocols

- ❑ A lock is a mechanism to control concurrent access to a data item
- ❑ Data items can be locked in two modes:
 1. exclusive (X) mode. Data item can be both reads as well as written. X-lock is requested using lock-X instruction.
 2. Shared (S) mode. Data item can only be read. S-lock is requested using lock-S instruction.
- ❑ Lock requests are made to concurrency-control manager. Transaction can proceed only after request is granted.



Lock-Based Protocols (Cont.)

- Lock-compatibility matrix

DATA ITEM (A)	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with lock already held on the item by other transactions
 -
 -
- Any number of transactions can hold shared locks on an item
- But if any transaction holds an exclusive lock on the item no other transaction may hold any lock on the item.

T_1	\bar{T}_2	T_1	\bar{T}_2	Lock Manager
<u>Read(A)</u>		- Lock-X(A)		GRANT-X(A, T_1)
$A = A - 100$		Read(A)		Release/ Revoke-X(A, T_1)
<u>Write(A)</u>		$A = A - 100$		
		Write(A)		
		Unlock-X(A)		
<u>Read(A)</u>			Lock-S(A)	GRANT-S(A, T_2)
<u>Read(C)</u>			Read(A)	Revoke-S(A, T_2)
<u>Read(B)</u>			Unlock-S(A)	
$B = B + 100$			Lock-S(C)	GRANT-S(C, T_2)
<u>Write(B)</u>			Read(C)	Revoke-S(C, T_2)
			Unlock-S(C)	GRANT-X(B, T_1)
		- Lock-X(B)		Revoke-X(B, T_1)
		Read(B)		
		$B = B + 100$		
		Write(B)		
		Unlock-X(B)		

Schedule with Lock Grants

- A locking protocol is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols enforce serializability by restricting the set of possible schedules.

T ₁	T ₂	Concurrency-control manager
<u>lock-X(B)</u> read(B) $B := B - 50$ write(B) <u>unlock(B)</u>		<u>grant-X(B, T₁)</u>
	<u>lock-S(A, T₂)</u> read(A) <u>unlock(A)</u> <u>lock-S(B)</u> read(B) <u>unlock(B)</u> <u>display(A+B)</u>	<u>grant-S(A, T₂)</u> <u>grant-S(B, T₂)</u>
<u>lock-X(A)</u> read(A) $A := A + 50$ write(A) <u>unlock(A)</u>		<u>grant-X(A, T₁)</u>

→ Revise - X(R, T_L)

Two Phase locking Protocol [2PL]

Lock & Unlock Request done in Two Phase

- ① Growing Phase (Acquire Lock)
- ② Shrinking Phase (Release Lock)

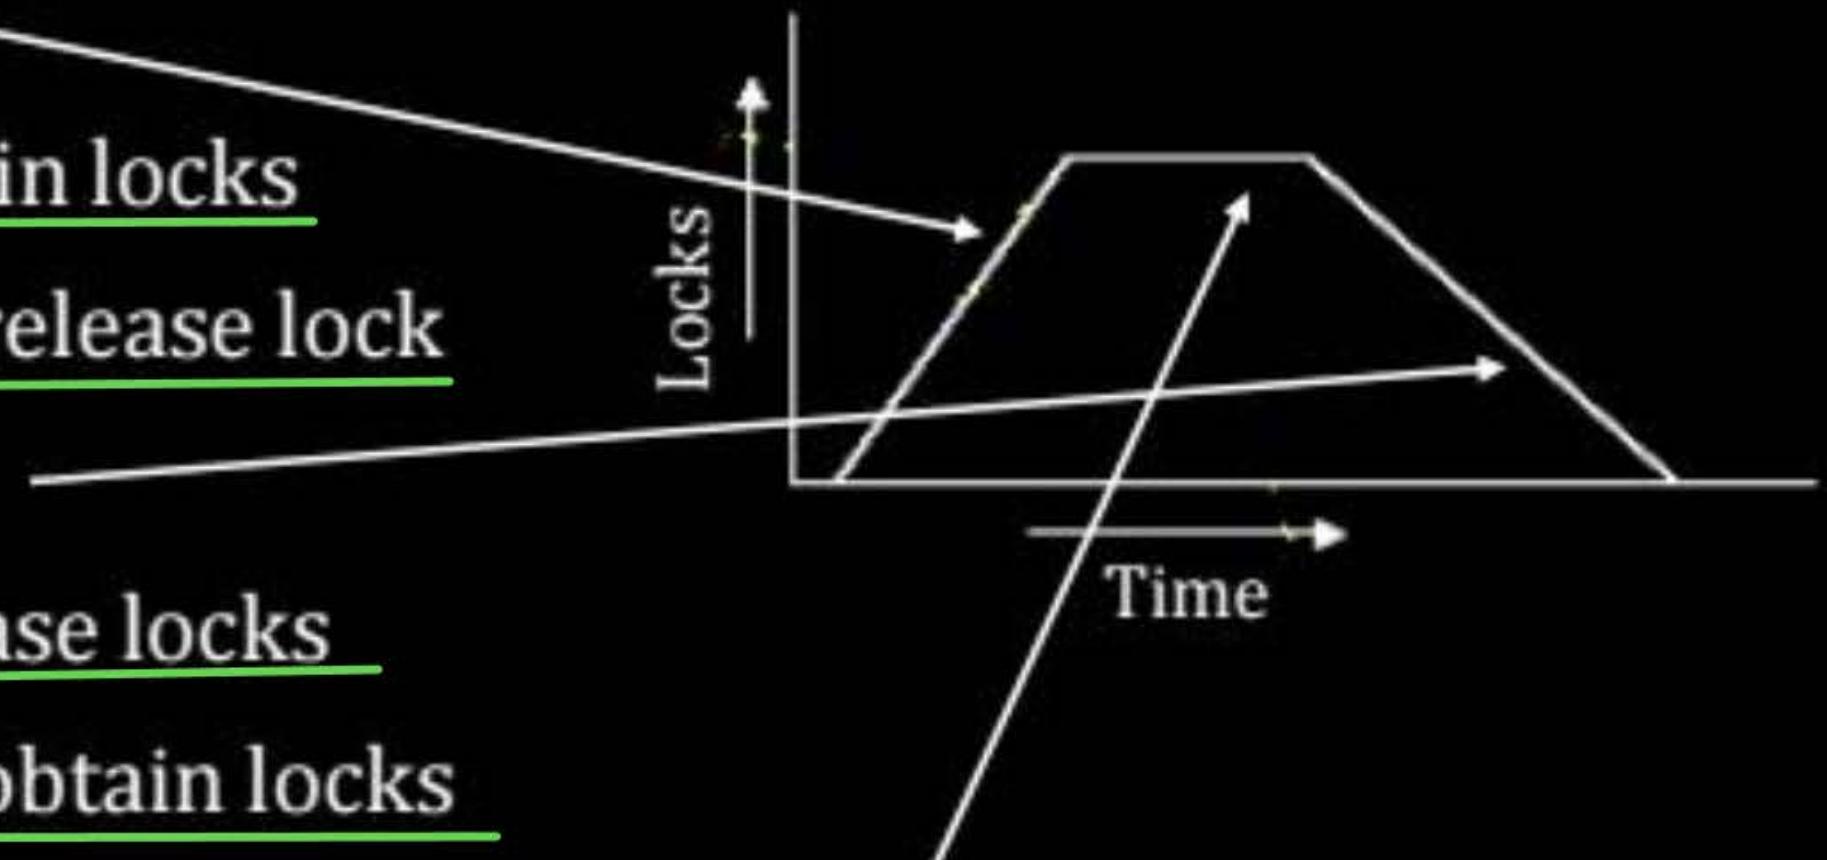
Note
Each Transaction first finish its Growing Phase then start Shrinking Phase.

① Gowning phase : In Gowning Phase, Transaction
May Obtain the Lock But
Must Not Release Any Lock .

② shrinking Phase : In Shrinking Phase,
Transaction Release the Lock
But Must Not ask for Any New Lock .
(Request)

The Two-Phase Locking Protocol

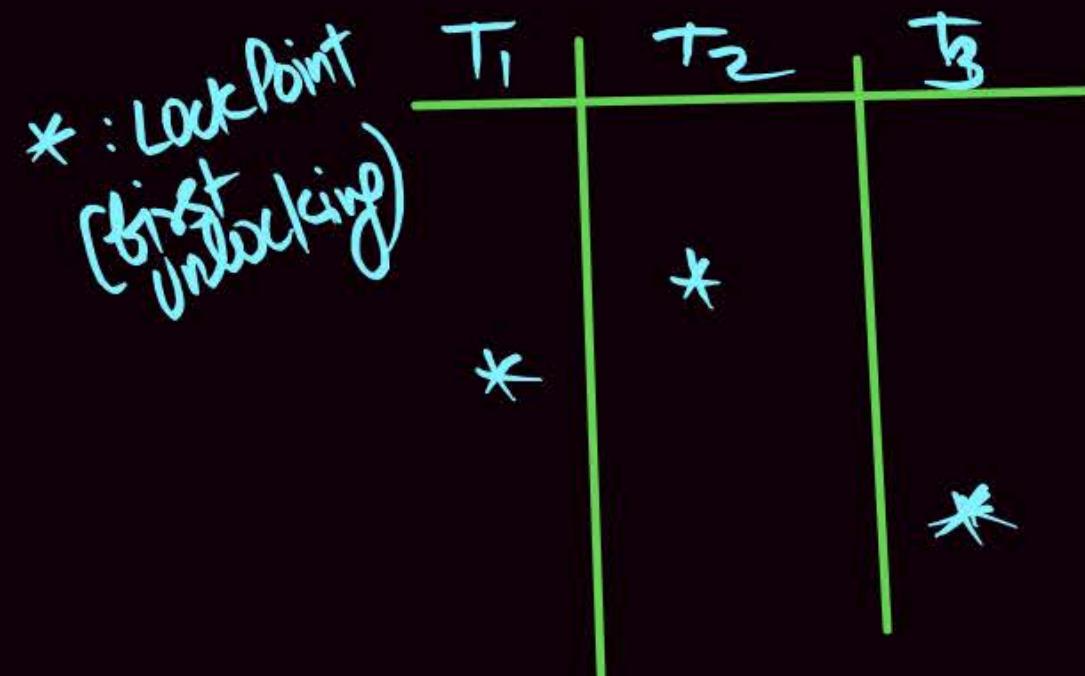
- A protocol which ensures conflict-serializable schedules.
- Phase 1: Growing Phase
 - ❖ Transaction may obtain locks
 - ❖ Transaction may not release lock
- Phase 2: Shrinking Phase
 - ❖ Transaction may release locks
 - ❖ Transaction may not obtain locks
- The protocol assures serializability. It can be proved that transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock).



LOCK POINT → Serializability Order.
LOCK POINT : Position of Last lock operation \textcircled{or}
First unlock operation.

\textcircled{OR}

Position from where Shrinking Phase of the transaction Start.



Serializability order.
 (T_2, T_1, T_3)

T_1	T_2	Lock Manager
$\overline{T_1}$	$\overline{T_2}$	\checkmark GRANT
Read(A) $A = A - 100$ <u>Write(A)</u>	LOCK-X(A) Read(A) $A = A - 100$ Write(A)	$\text{LOCK-S}(A)$ $\times \times$ No
Read(A)	① Lock-X(B) ② Unlock-X(A)	\rightarrow Release-X(A, T_1)
Read(C)	LOCK-S(A) Read(A)	\checkmark YES(GRANT)
Read(B) $B = B + 100$ <u>Write(B)</u>	③ LP ④ LOCK-S(C)	\checkmark YES.
	Unlock-S(A) Read(C) Unlock-S(C)	\rightarrow ✓ (T_1, T_2) .
	Read(B) $B = B + 100$ <u>Write(B)</u> Unlock-X(B)	\checkmark

Important Point about 2PL.

- ① 2PL Ensure Conflict Serializability (Serializability)
If a Schedule is allowed (followed) by 2PL then
its Ensure Conflict Serializable Schedule (Serializability).
- ② Serializability Order determined by Lock Point

Important Point about 2PL.

③ 2PL Not Ensure Recoverability.

Irrecoverable Schedule followed by 2PL.

T_1	T_2
$R(A)$	
$W(A)$	
$R(A)$	
Commit	
Commit	
Irrecoverable Schedule	

T_1	T_2
$X(A)$	
$R(A)$	
$W(A)$	
unlock-X(A)	
$S(A)$	
$R(A)$	
unlock-S(A)	
commit	
commit	

Irrecoverable Schedule
But allowed
(followed) by QPL.

Important Point about 2PL.

④ 2PL Suffers from (May Not Free From) Deadlock.

①

T_1	T_2
T_1	X(A) W(A)
T_2	S(B) R(B)
$w(B)$	$R(A)$
$R(A)$	$X(B)$

Denied by $T_2 \rightarrow X(B)$

$S(A) \leftarrow$ Denied by T_1

Deadlock

②

T_1	T_2
T_1	R(A)
T_2	W(B)
$R(B)$	$R(A)$
$W(A)$	$S(B)$

Denied by $T_2 \rightarrow S(B)$

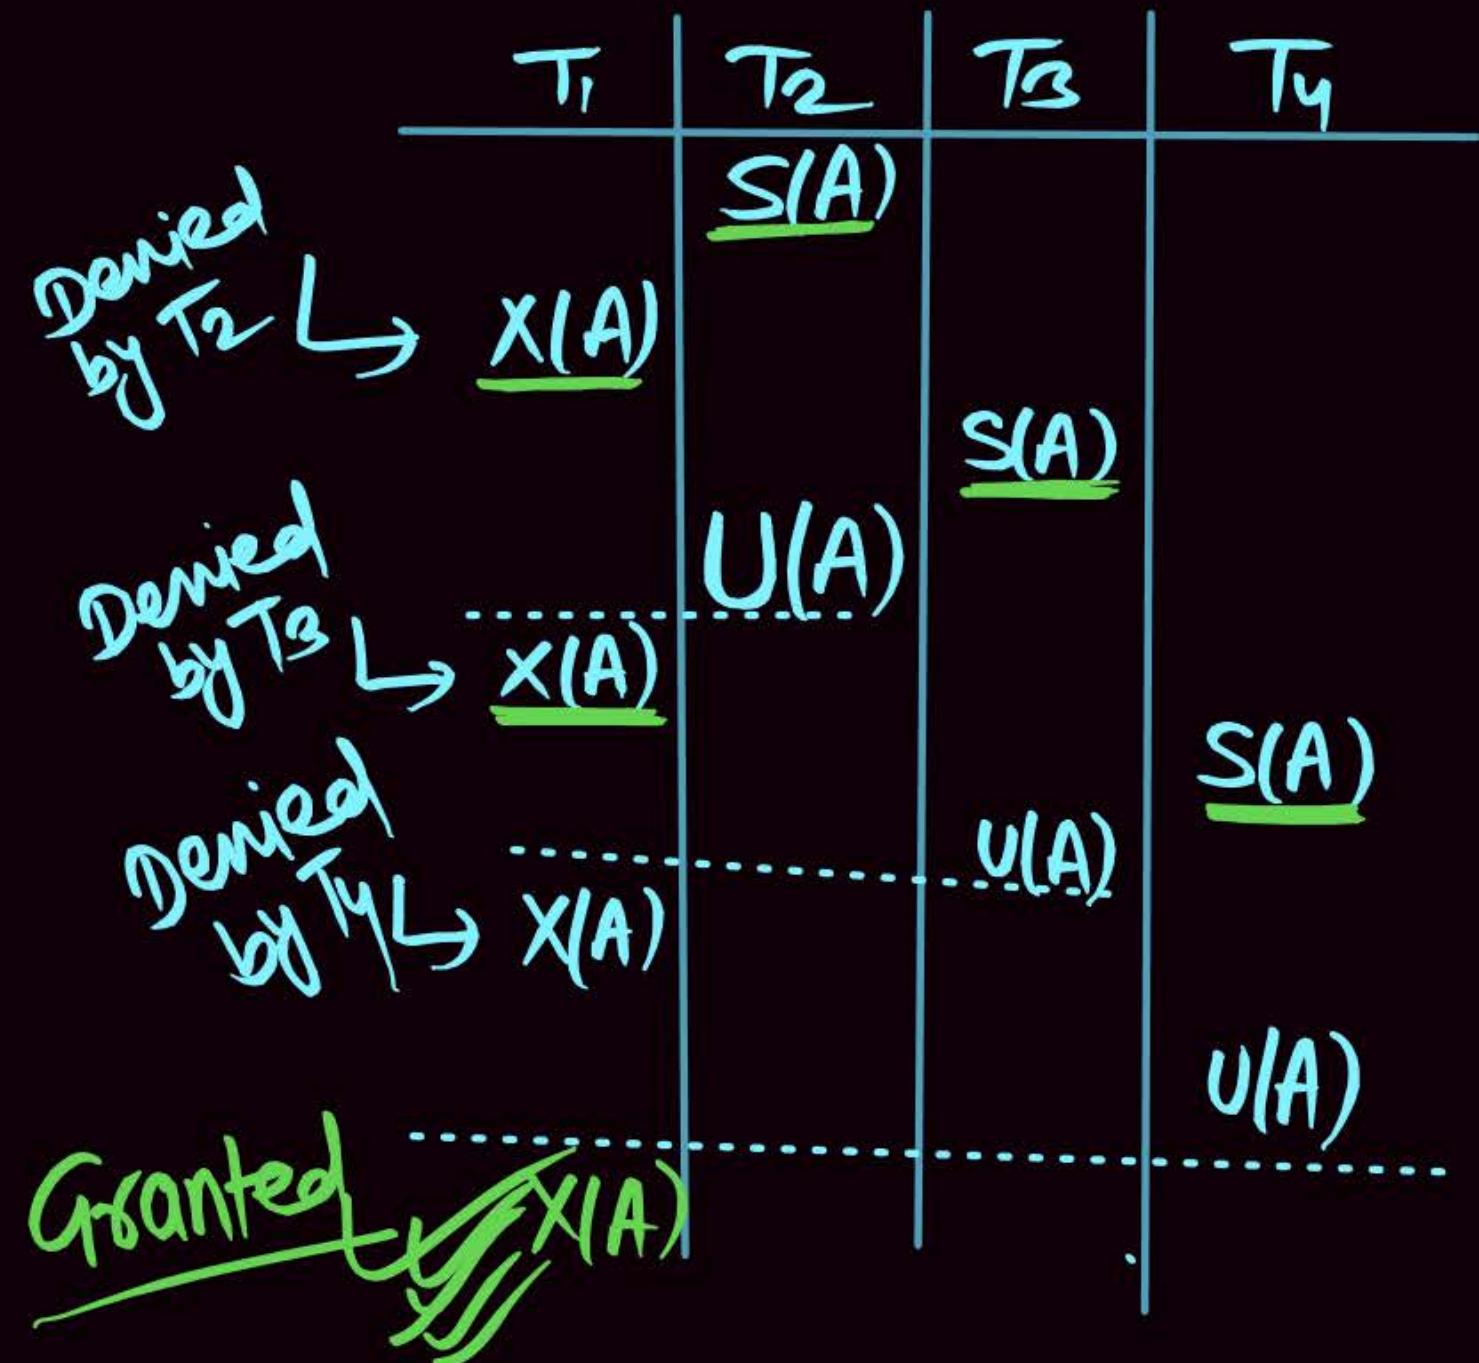
$X(B) \leftarrow$ Denied by T_1

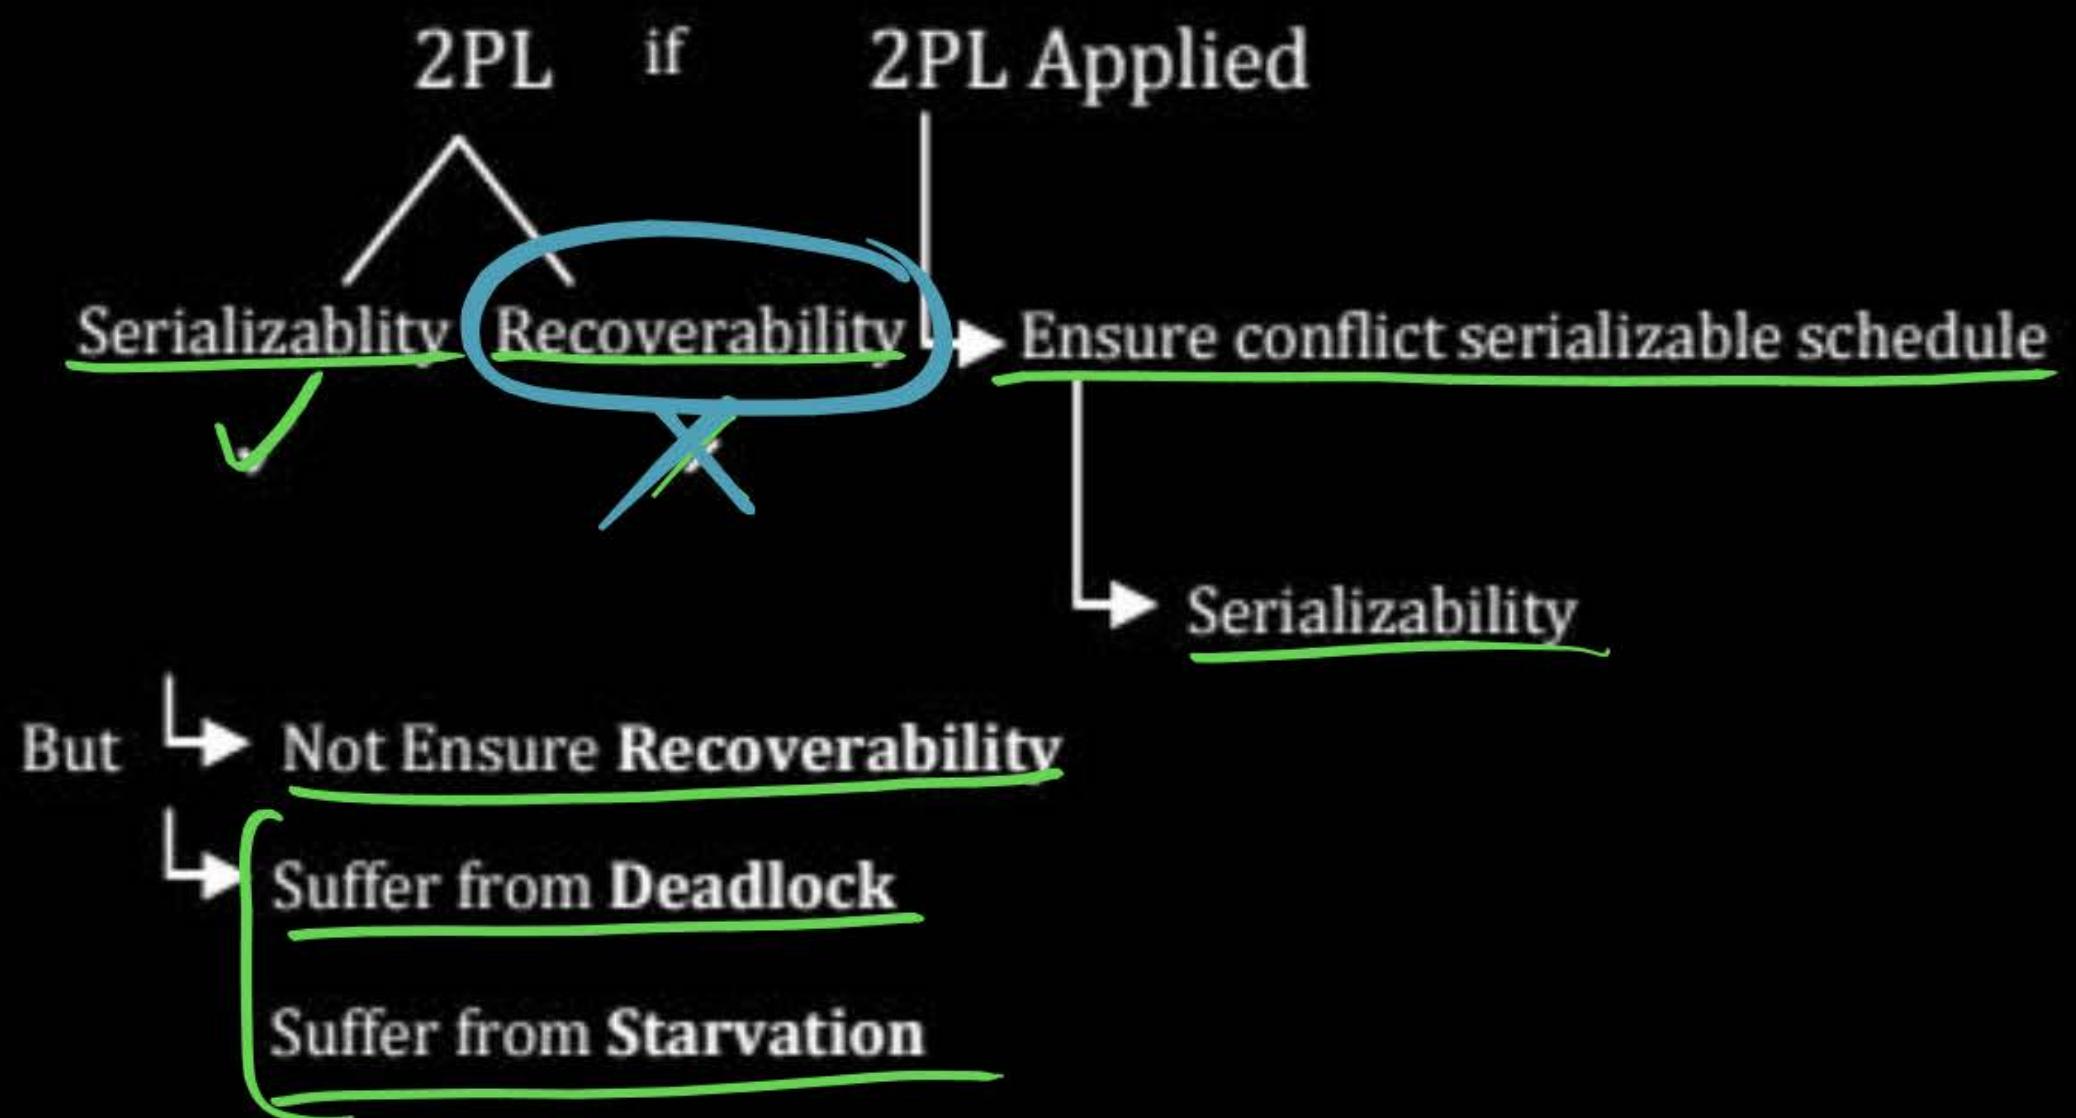
$X(A) \leftarrow$ Denied by T_1

Deadlock

Important Point about 2PL.

⑤ 2PL Subject from Starvation.





STRICT 2PL : 2PL + ALL Exclusive ^(X Lock) Lock

taken by the transaction

Must be Held until Commit/Rollback

OR X Lock Release After Commit/Rollback

T _i
S(A)
X(B)
X(C)
S(D)
=
U(A)
U(D)
C R
U(B)
U(C)



Strict Recoverable

- Conflict Serializable
- Recoverable Schedule
- Cascaded (No Cascading Rollback).
- Strict Recoverable

Suffer from

- Deadlock
- Starvation

Rigorous 2PL: 2PL + All Locks [Shared [S] & Exclusive [X]]
Must be Held by the transaction

until Commit / Rollback

OR

2PL + ALL lock [S & X] Release

After Commit

Rigorous 2PL

↳ Suffer from Deadlock

↳ Starvation.

T _i
S(A)
X(B)
X(C)
S(D)
=
C R
U(A)
U(B)
U(C)
U(D)

The Two-Phase Locking Protocol (Cont.)

- Two-phase locking does not ensure freedom from deadlocks
- Extensions to basic two-phase locking needed to ensure recoverability of freedom from cascading roll-back
 - ❖ Strict two-phase locking: a transaction must hold all its exclusive locks till it commits/aborts.
 - Ensures recoverability and avoids cascading roll-backs
 - ❖ Rigorous two-phase locking: a transaction must hold all locks till commit/abort.
 - Transactions can be serialized in the order in which they commit.
 - Most databases implement rigorous two-phase locking, but refer to it as simply two-phase locking

Conservative 2PL ∵ Each transaction Acquire

ALL the Lock in the beginning(at ^{the} start) of Execution & Release After Commit.

Ensure

- Conflict Serializability
- Recoverability
- Cascadels
- No Cascading Rollback
- Strict Recoverable
- No Deadlock (Freedom From Deadlock)

But suffer from
Starvation.

P
W

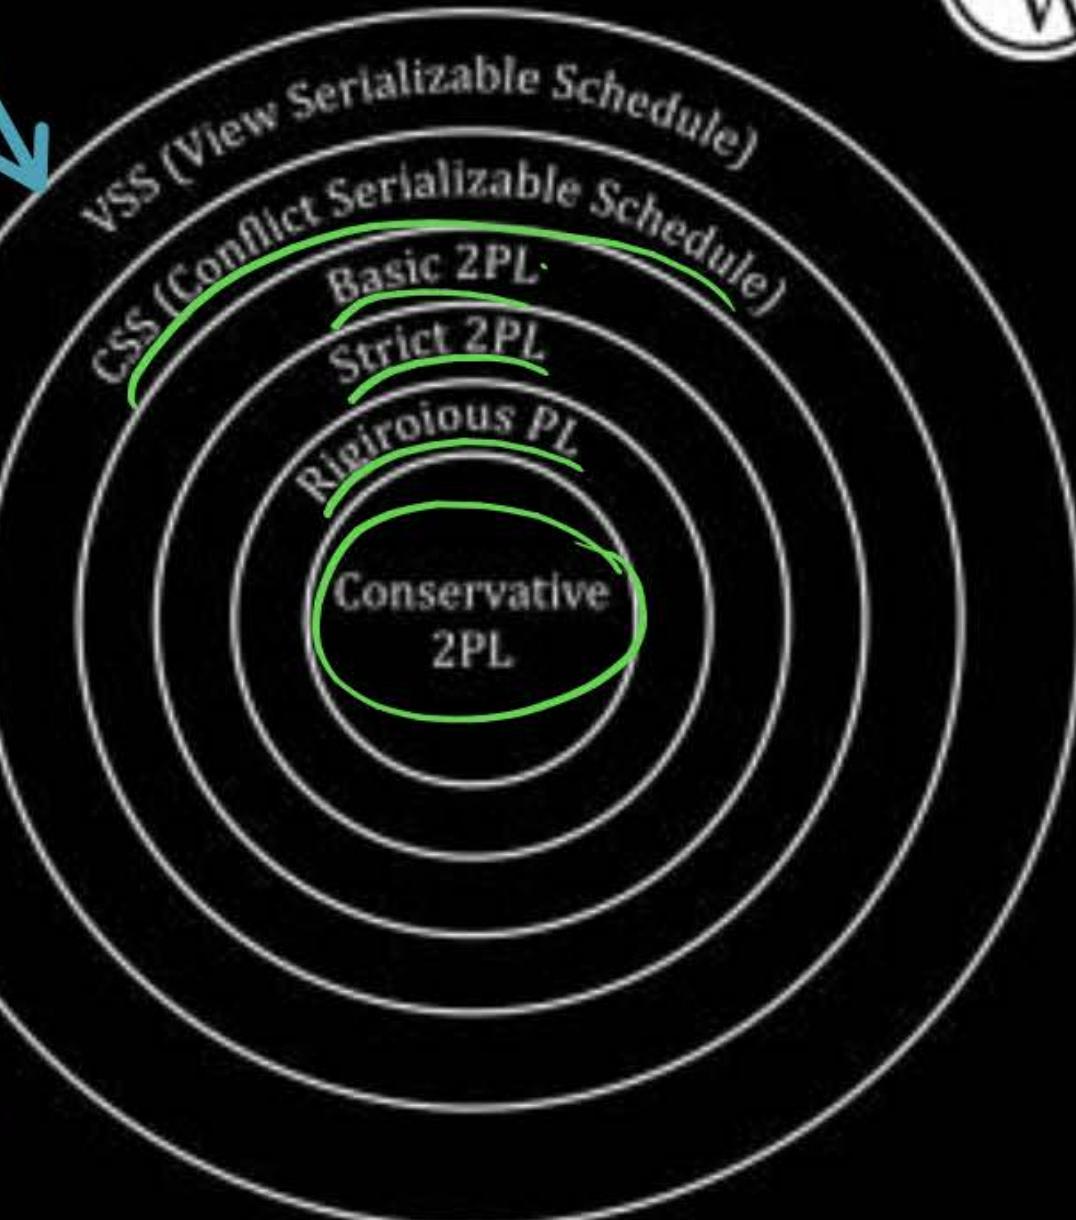
τ_1	τ_1	τ_1	τ_1	τ_1
1	2	3	4	5
Lock-S(A)	Lock-S(A)	Lock-S(A)	Lock-S(A)	Lock-S(A)
R(A)	R(A)	R(A)	R(A)	R(A)
Lock-X(B)	Lock-X(B)	Lock-X(B)	Lock-X(B)	Lock-X(B)
R(B)				R(B)
Unlock(A)	Unlock(A)			
	R(B)	R(B)	R(B)	
	W(B)	W(B)	W(B)	
Unlock(B)	Commit	Commit	Commit	Commit
Commit	Unlock(B)	Unlock(A)	Unlock(A)	Unlock(B)
				Unlock(B)
				Commit

Basic 2PL

Strict 2PL
(Basic 2PL)

Rigorous 2PL
Strict & Basic 2PL

Conservative 2PL
Rigorous, Strict & Basic 2PL



TIMESTAMP BASED CONCURRENCY CONTROL

Timestamp-Based Protocols

- Each transaction T_i is issued a timestamp $TS(T_i)$ when it enters the system.
 - ❖ Each transaction has a unique timestamp
 - ❖ Newer transaction have timestamp strictly greater than earlier ones
 - ❖ Timestamp could be based on a logical counter
 - Real time may not be unique
 - Can use (wall-clock time, logical counter) to ensure
- Timestamp-based protocols manage concurrent execution such that **time-stamp order = serializability order**
- Several alternative protocols based on timestamps

Timestamp-Based Protocols

The **timestamp ordering (TSQ) protocol**

- Maintains for each data Q two timestamp values:
 - ❖ **W-timestamp(Q)** is the largest time-stamp of any transaction that executed **write** (Q) successfully.
 - ❖ **R-timestamp (Q)** is the largest time-stamp of any transaction that executed **read** (Q) successfully.
- Imposes rules on read and write operations to ensure that
 - ❖ any conflicting operations are executed in timestamp order
 - ❖ out of order operations cause transaction rollback

I: T_i - Read(Q) (Transaction T_i Issue R(Q) Operation)

- (i) If $TS(T_i) < WTS(Q)$: Read operation Reject & T_i Rollback.
- (ii) If $TS(T_i) \geq WTS(Q)$: Read operation is allowed
and Set Read - TS(Q) = $\max[RTS(Q), TS(T_i)]$

II: T_i - Write(Q) (Transaction T_i Issue Write(Q) Operation)

- (i) If $TS(T_i) < RTS(Q)$: Write operation Reject & T_i Rollback.
- (ii) If $TS(T_i) < WTS(Q)$: Write operation Reject & T_i Rollback.
- (iii) Otherwise execute write (Q) operation
Set Read WTS(Q) = $TS(T_i)$

Timestamp-Based Protocols (Cont.)

- Suppose a transaction T_i issues a **read** (Q)
 1. If $TS(T_i) \leq W\text{-timestamp } (Q)$, then T_i needs to read a value of Q that was already overwritten.
 - Hence, the **read** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) \geq W\text{-timestamp } (Q)$, then the **read** operation is executed, and $R\text{-timestamp}(Q)$ is set to.
 $\max(R\text{-timestamp } (Q), TS (T_i))$.

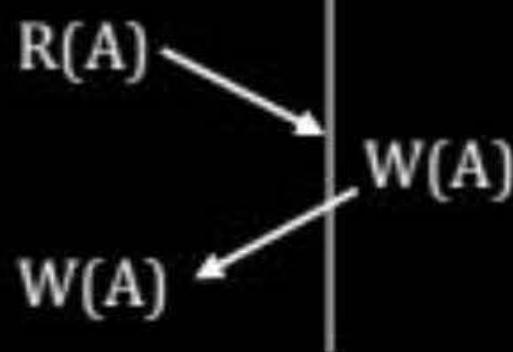
Timestamp-Based Protocols (Cont.)

- Suppose a transaction T_i issues **write(Q)**
 1. If $TS(T_i) < R\text{-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that the value would never be produced.
 - Hence, the **write** operation is rejected, and T_i is rolled back.
 2. If $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q.
 - Hence, this **write** operation is rejected, and T_i is rolled back.
 3. Otherwise, the **write** operation is executed, and $W\text{-timestamp}(Q)$ is set to $TS(T_i)$.

P
W

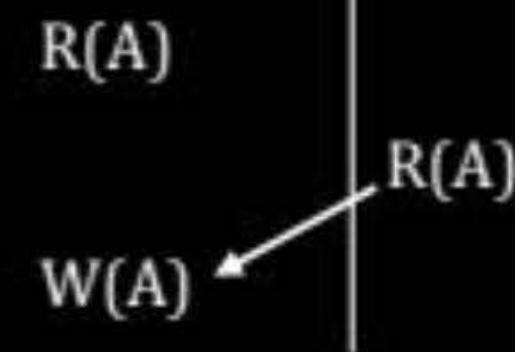
(1)

T ₁ (10)	T ₂ (20)
---------------------	---------------------



(2)

T ₁ (10)	T ₂ (20)
---------------------	---------------------



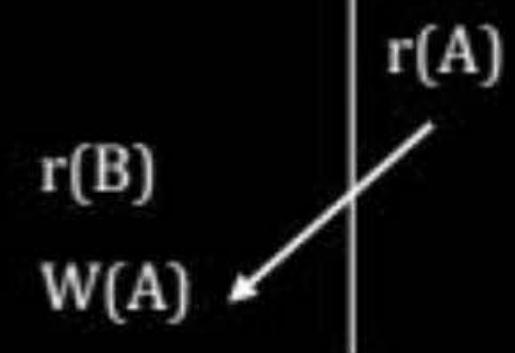
(3)

T ₁ (10)	T ₂ (20)
---------------------	---------------------



(4)

T ₁ (10)	T ₂ (20)
---------------------	---------------------



TS (T₁) < TS (T₂)

T₁ → T₂

TS (T₁) < TS (T₂)

T₁ → T₂

Thomas' Write Rule

- ❑ Modified version of the timestamp-ordering protocol in which obsolete **write** operations may be ignored under certain circumstances.
- ❑ When T_i attempts to write data item Q , if $TS(T_i) < W\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of $\{Q\}$.
 - ❖ Rather than rolling back T_i as the timestamp ordering protocol would have done, this **{write}** operation can be ignored.
- ❑ Otherwise this protocol is the same as the timestamp ordering protocol.
- ❑ Thomas' Write Rule allows greater potential concurrency.
 - ❖ Allows some view-serializable schedules that are not conflict-serializable.

Thomas Write Rule (View Serializability)

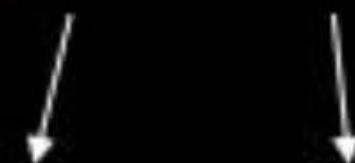
1. $TS(T_i) < RTS(Q)$: Rollback
2. $TS(T_i) < WTS(Q)$: Write operation is Ignored and No Roll back

Same as TSP

Time Stamp Protocol: Ensure serializability deadlock free but starvation possible

Deadlock Prevention Algorithm

(1) Wait-Die

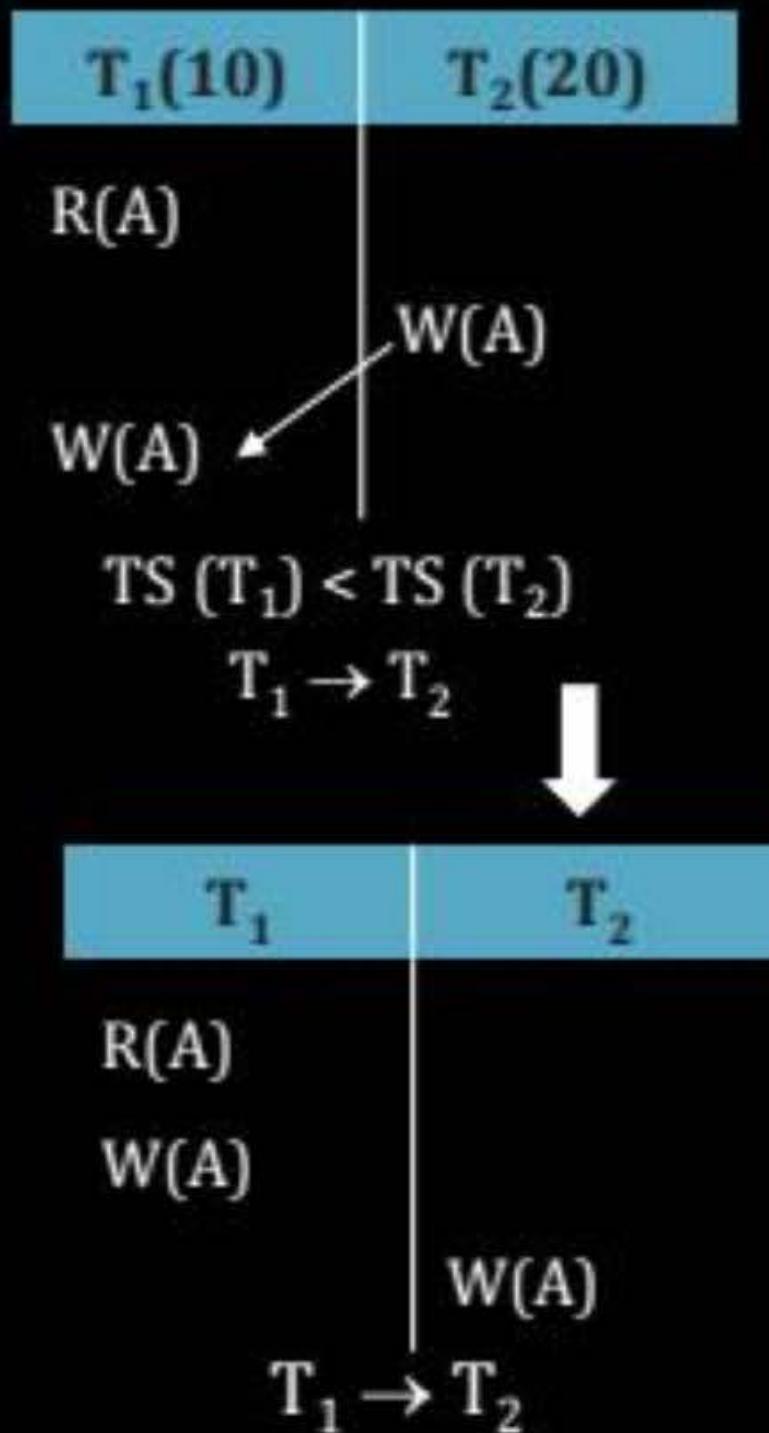


Older

(2) Wound-wait



Younger



Q.

In which one of the following Lock Scheme Deadlock cannot occur?

[MCQ]

P
W

- A** Basic 2PL
- B** Strict 2PL
- C** Conservative 2PL
- D** Rigorous 2PL

Q.

Consider the following statement about lock-based protocol

P
W

- (A) 2 PL (2phase locking) protocol Ensure view serializability
 - (B) 2PL ensure recoverability & No cascading rollback.
 - (C) Strict 2 PL ensure recoverability & no cascading rollback.
 - (D) Strict 2 PL avoids deadlock (not suffering from deadlock).
- How many numbers of above statement are correct?

A

1

[MCQ]

B

2

C

3

D

4

Q.

Consider the following Schedule:

[MSQ]

P
W

$r_1(x) \ r_2(y) \ r_2(x) \ w_1(z) \ r_1(y) \ w_3(y) \ r_3(z) \ w_2(y) \ w_3(x)$

which of the following time stamp ordering Not allows to execute schedule using Thomas Write rule time stamp Ordering Protocol?

- A** $(T_1, T_2, T_3) = (20, 30, 10)$
- B** $(T_1, T_2, T_3) = (10, 20, 30)$
- C** $(T_1, T_2, T_3) = (10, 30, 20)$
- D** $(T_1, T_2, T_3) = (30, 20, 10)$

Consider the following two statements about database transaction schedules:

- I. Strict two-phase locking protocol generates conflict serializable schedules that are also recoverable.
- II. Timestamp-ordering concurrency control protocol with Thomas Write Rule can generate view serializable schedules that are not conflict serializable.

Which of the above statements is/are True?

[GATE-2019-CS: 1M]

A I only

B II only

C Both I and II

D Neither I nor II

Consider the following partial Schedule S involving two transactions T1 and T2. Only the read and the write operations have been shown. The read operation on data item P is denoted by read (P) and the write operation on data item P is denoted by write (P).

Suppose that the transaction T1 fails immediately after time instance 9. Which one of the following statements is correct?

Time	Transaction-id	
	T 1	T 2
1.	read(A)	
2.	write(A)	
3.		read(C)
4.		write(C)
5.		read(B)
6.		write(B)
7.		read(A)
8.		commit
9.	read(B)	

[GATE-2015-CS: 2M]

- A T2 must be aborted and then both T1 and T2 must be restarted to ensure transaction atomicity
- B Schedule S is non recoverable and cannot ensure transaction atomicity
- C One T2 must be aborted and then restarted to ensure transaction atomicity
- D Schedule S is recoverable and can ensure atomicity and nothing else needs to be done

Which of the following concurrency control protocols ensure both conflict serializability and freedom from deadlock?

- I. 2-phase locking
- II. Time-stamp ordering

[GATE-2010-CS: 1M]

A I only

B II only

C Both I and II

D Neither I nor II

MCQ

Consider a simple checkpointing protocol and the following set of operations in the log.

{start, T4}; {write, T4, y, 2, 3}; {start, T1};

{commit, T4}; {write, T1, z, 5, 7};

{checkpoint}:

{start, T2}; {write, T2, x, 1, 9}; {commit, T2};

{start, T3}; {write, T3, z, 7, 2};

If a crash happens now and the system tries to recover using both undo and redo operations, what are the contents of the undo list and the redo list?

After Check Point [GATE-2015-CS: 2M]

Commit T₂: Redo T₂

Undo: T₃ & T₁.

- A Undo: T₃, T₁; Redo: T₂
- B Undo: T₃, T₁; Redo: T₂, T₄
- C Undo: none; Redo: T₂, T₄, T₃, T₁
- D Undo: T₃, T₁, T₄; Redo: T₂

Those Transaction Commit before Check Point Not Required
Any Redo operation.

Those Transaction Commit After check point Require Redo operation

Those Transaction Not Commit Require Undo operation.

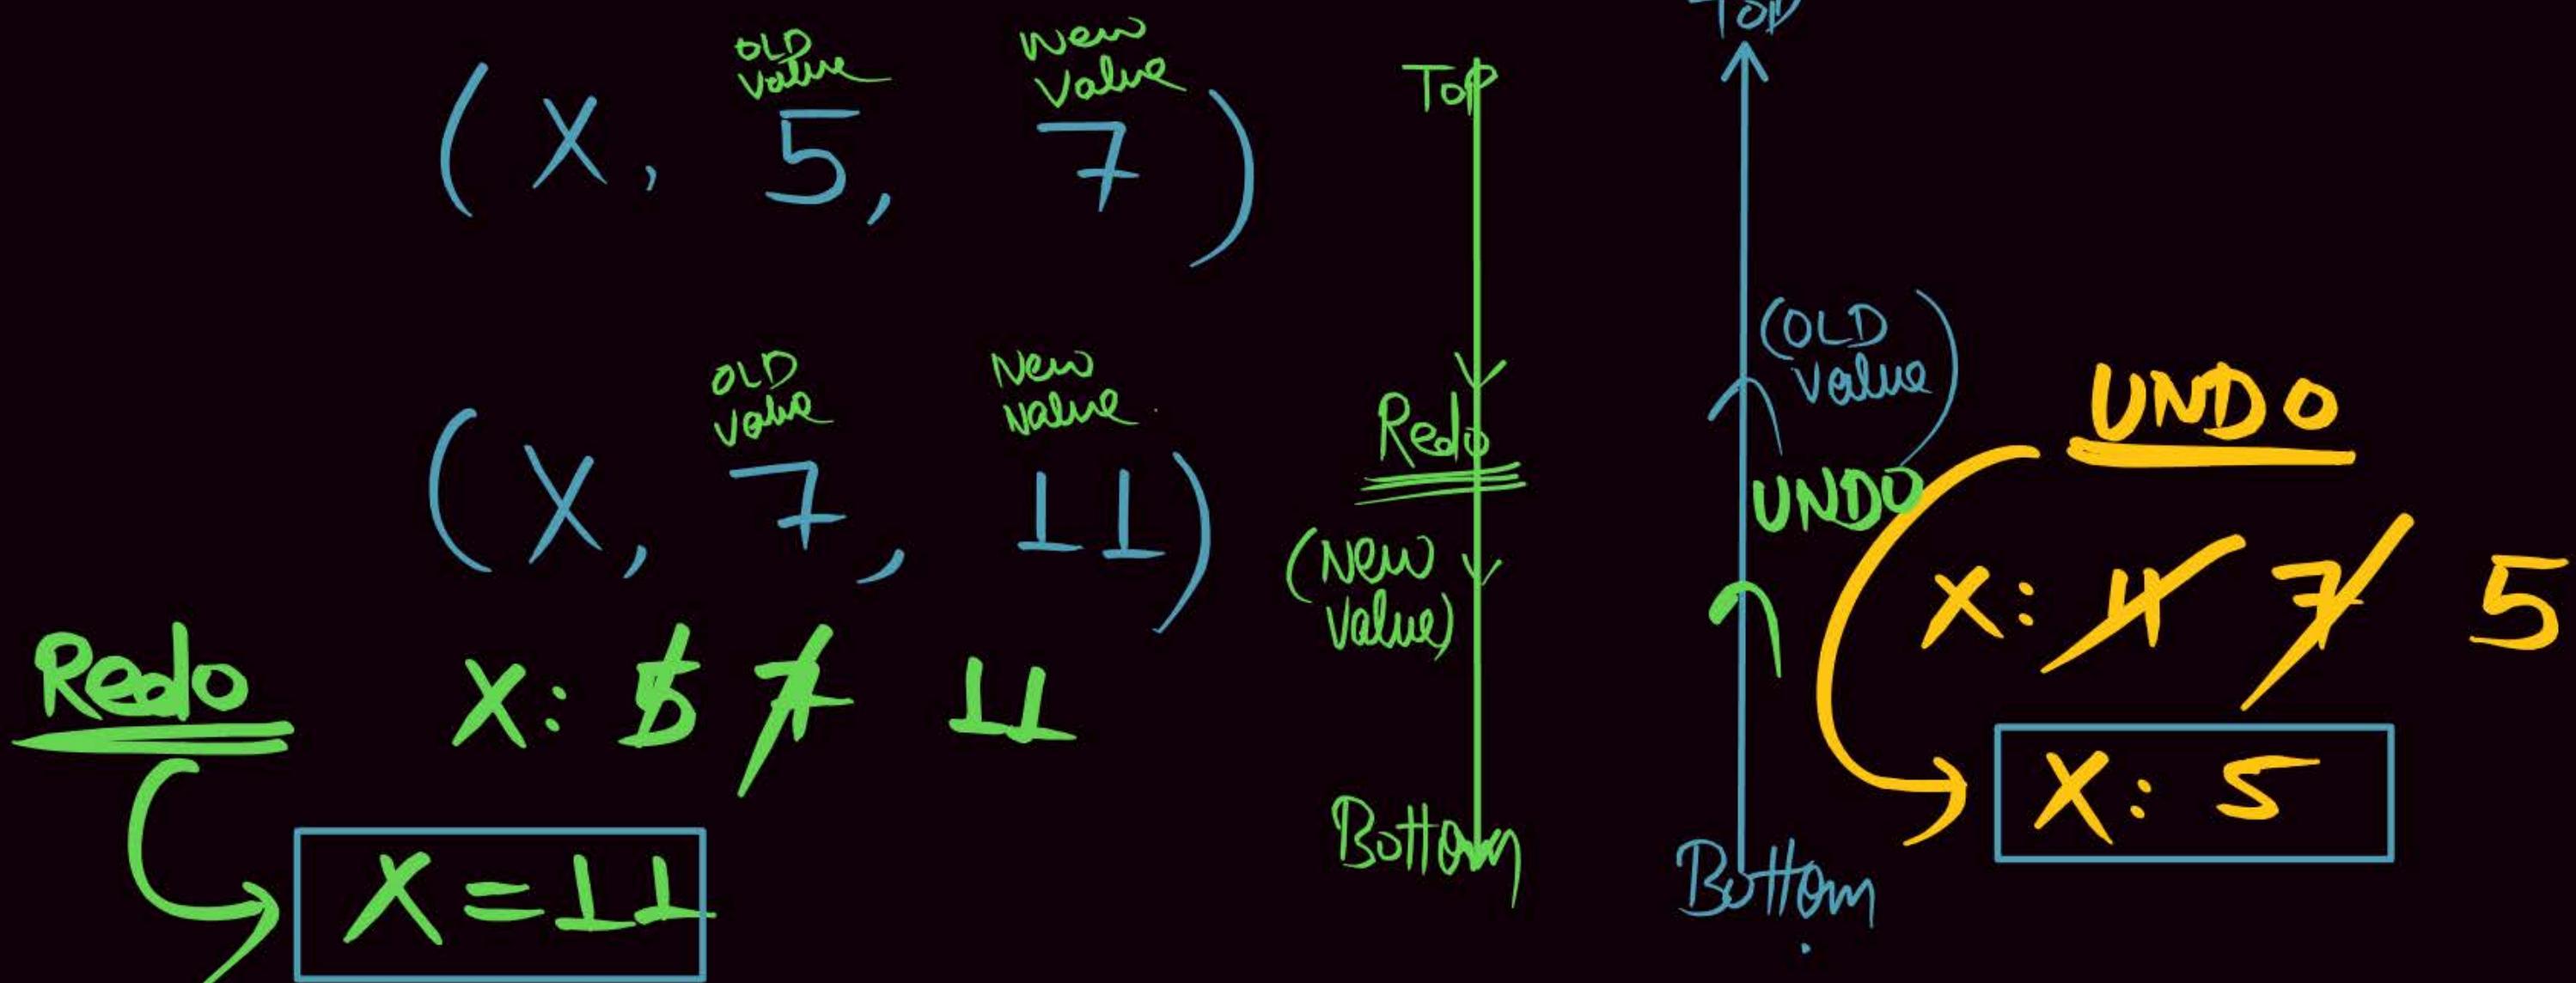
If T_i Commit then Redo.

If Any Transaction that Not Commit Require UNDO
operation.

Log Based Recovery

Redo & Undo Concept

Transaction number (observation, Data Item, ^(old) OLD value, ^(new) New value)



Q.

Consider the following database schedule with two transactions, T_1 and T_2 .

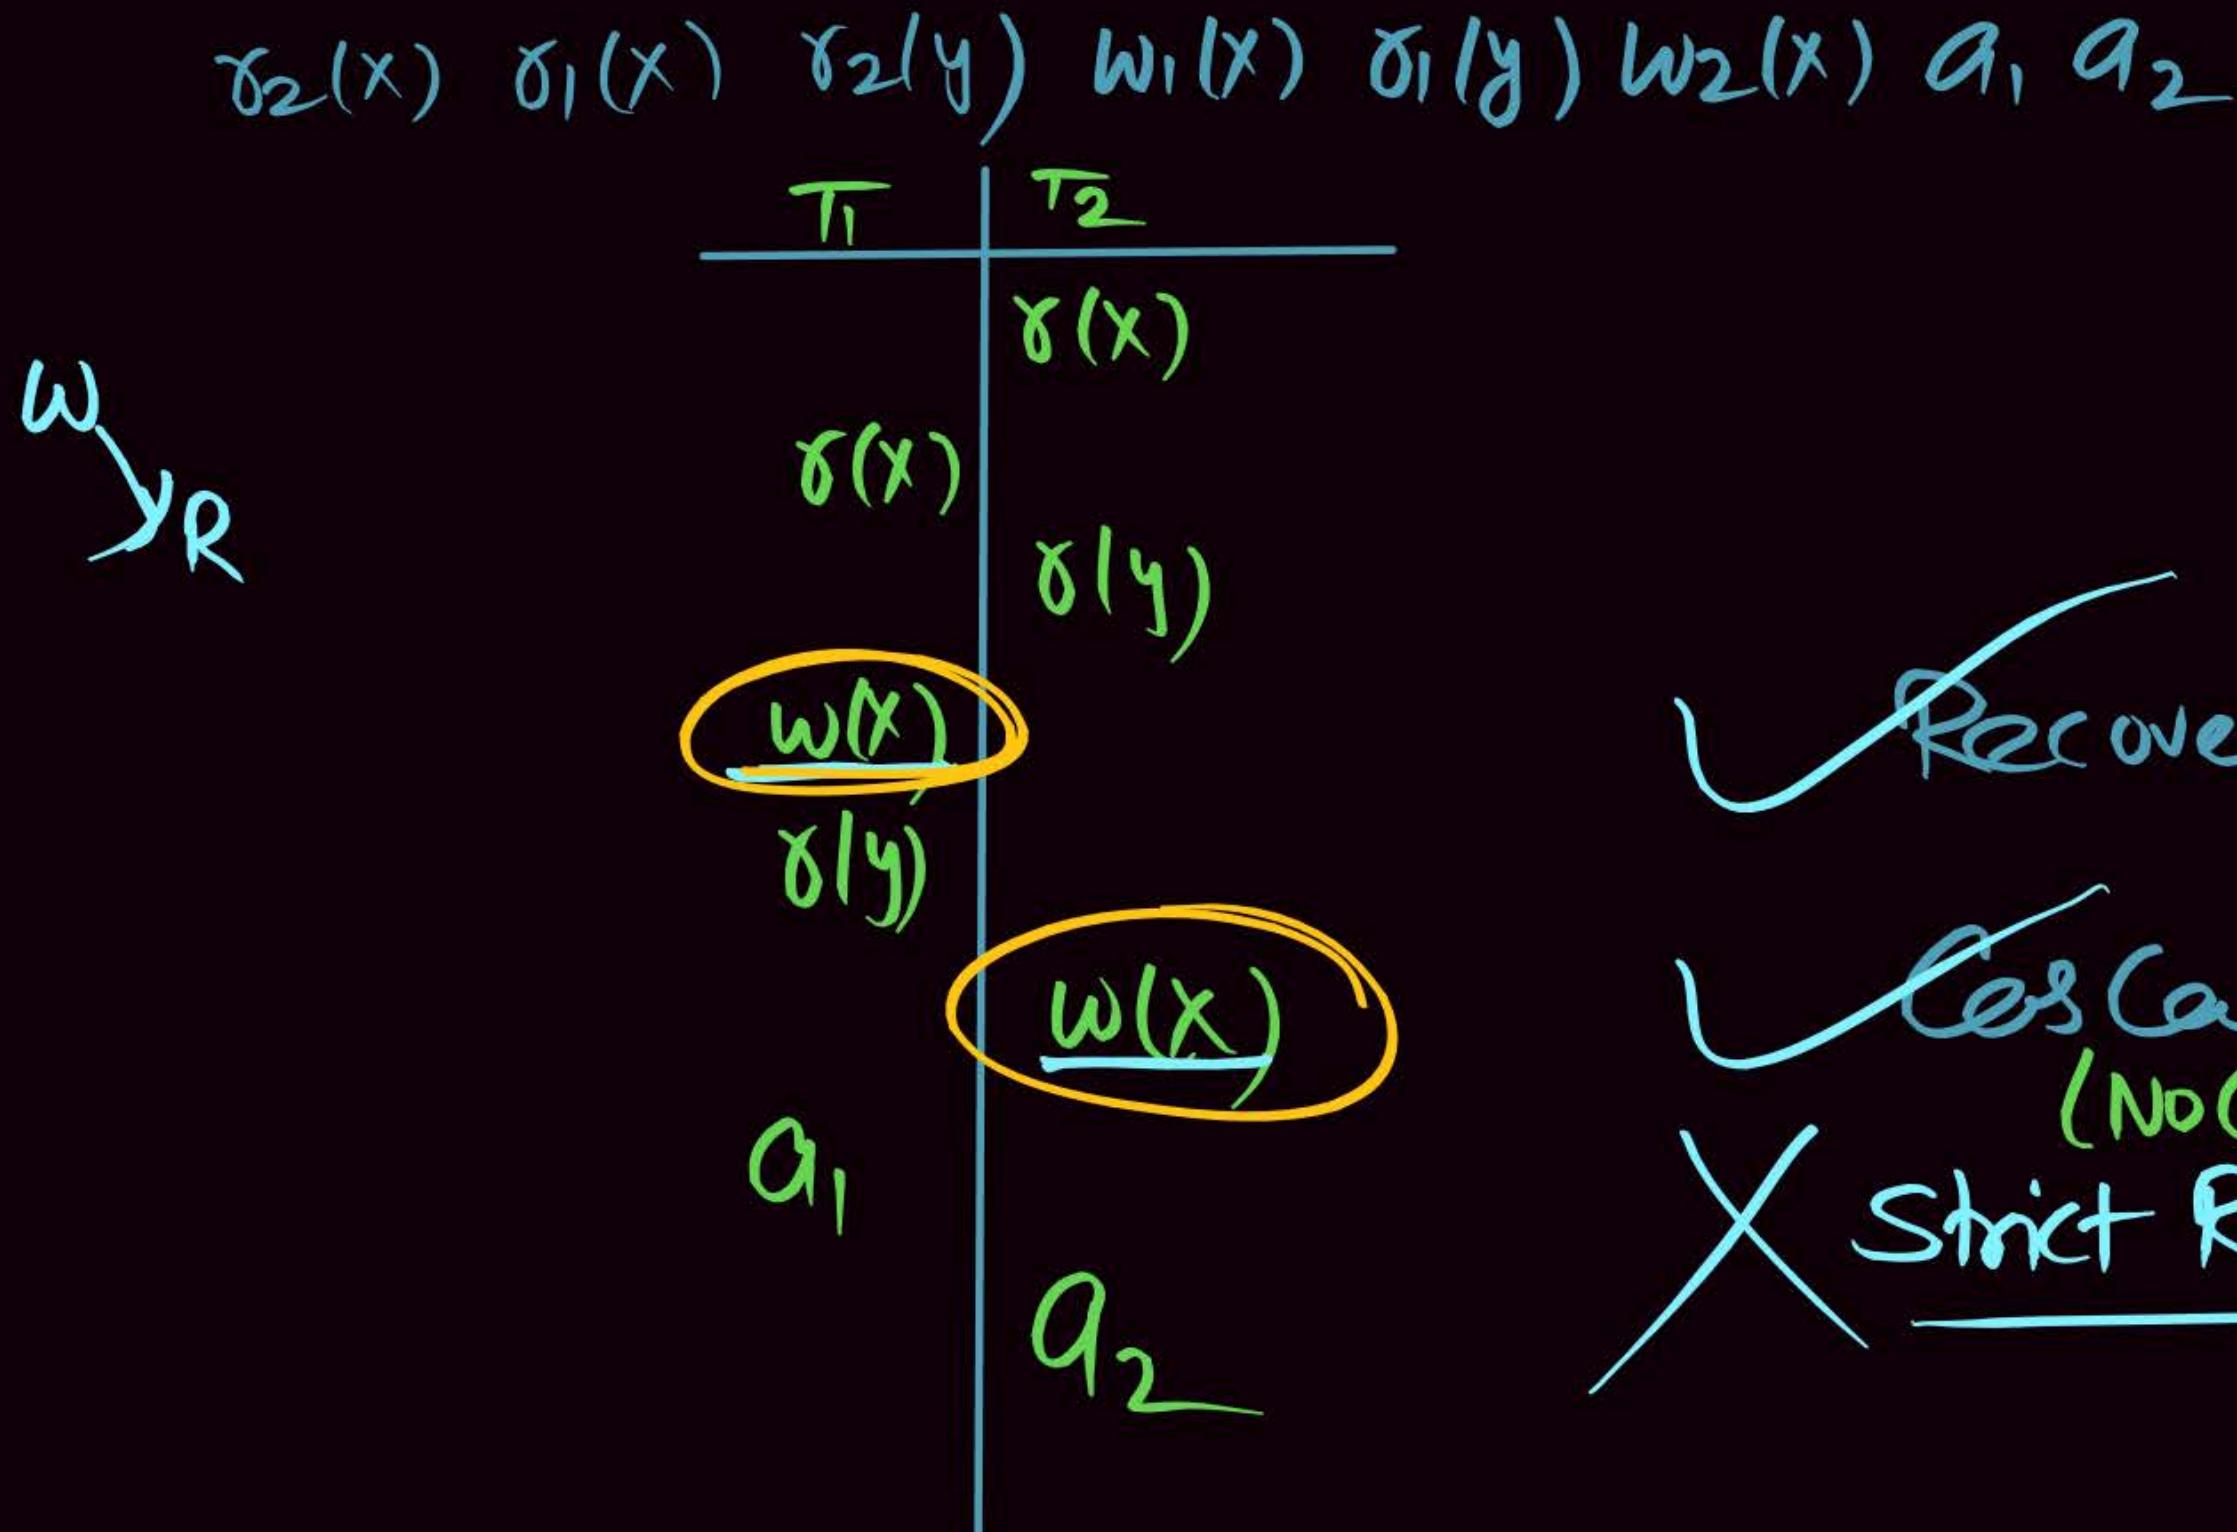
$$S = r_2(X); r_1(X); r_2(Y); w_1(X); r_1(Y); w_2(X); a_1; a_2$$

where $r_i(Z)$ denotes a read operation by transaction T_i on a variable Z , $w_i(Z)$ denotes a write operation by T_i on a variable Z and a_i denotes an abort by transaction T_i

Which one of the following statements about the above schedule is TRUE?

[MCQ:2016–2M]

- A S is non-recoverable
- B S is recoverable, but has a cascading abort
- C S does not have a cascading abort
- D S is strict



~~✓ Recoverable~~
~~✓ Cascaded~~
~~(No Cascading Rollback)~~
~~✗ Strict Recoverable~~

Q.

Let S be the following schedule of operations of three transactions T_1 , T_2 and T_3 in a relational database system:

$R_2(Y), R_1(X), R_3(Z), R_1(Y), W_1(X), R_2(Z), W_2(Y), R_3(X), W_3(Z)$

Consider the statements P and Q below:

P: S is conflict-serializable.

Q: If T_3 commits before T_1 finishes, then S is recoverable.

Which one of the following choices is correct?

A Both P and Q are true.

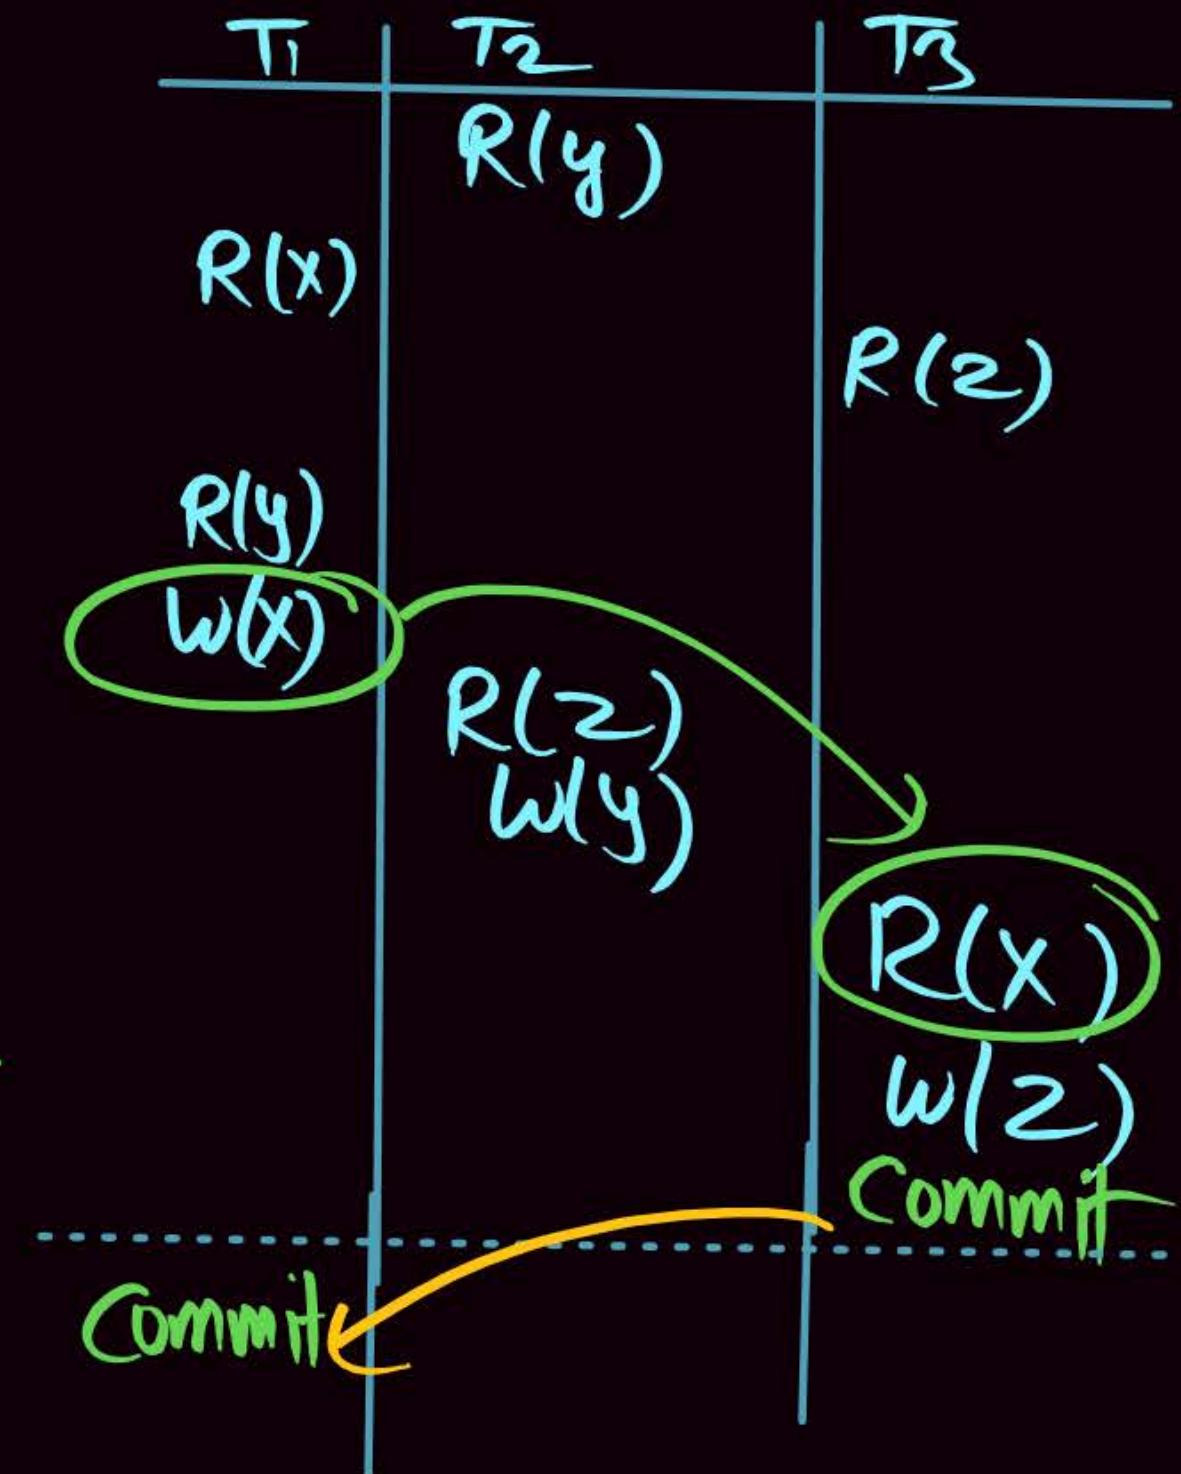
[MCQ: 2021-2M]

B P is true and Q is false.

C P is false and Q is true.

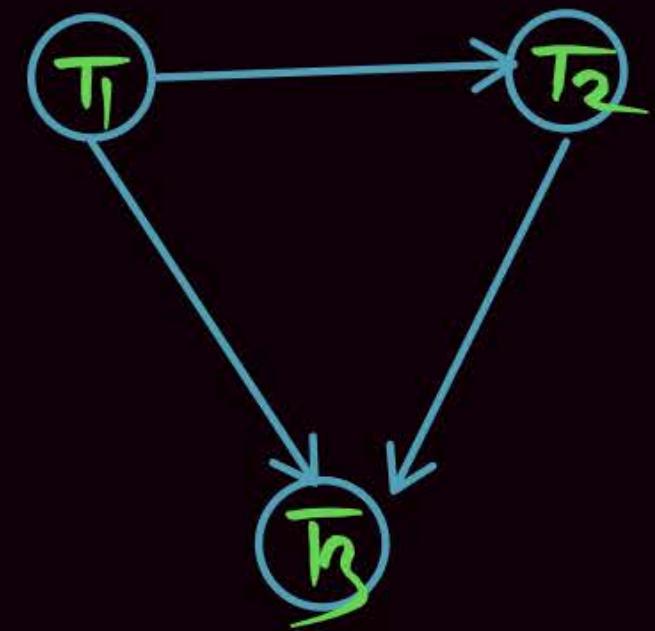
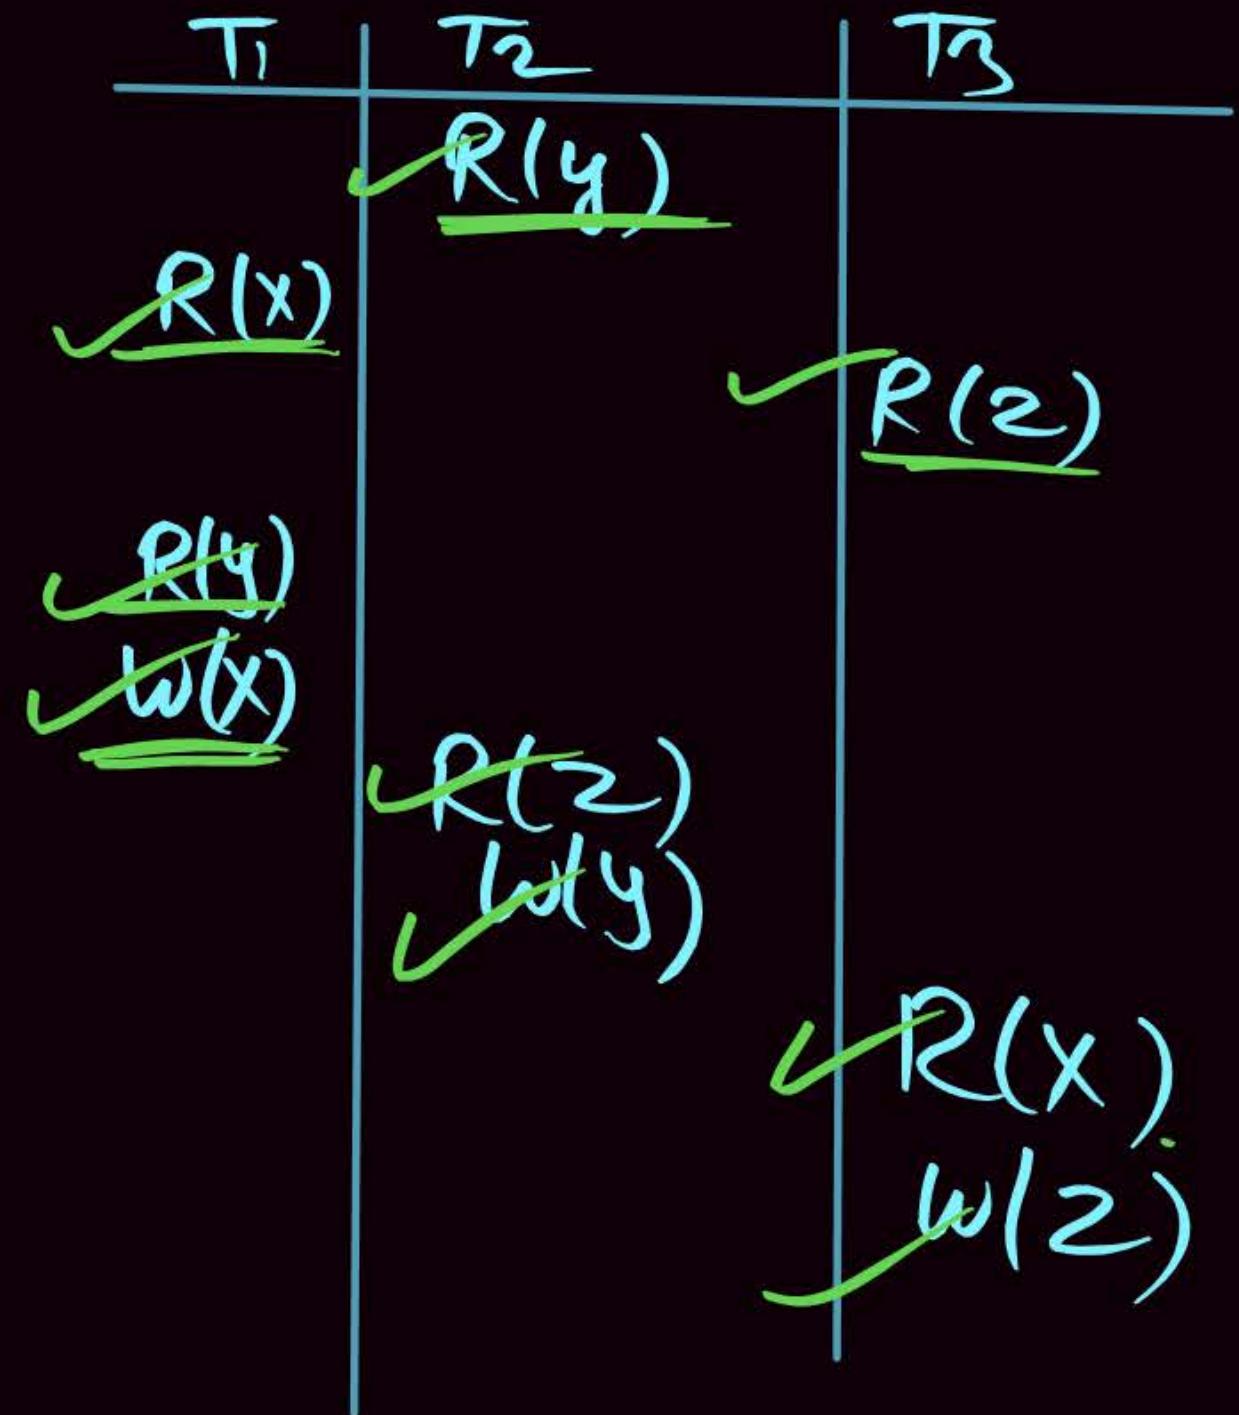
D Both P and Q are false.

$R_2(y)$ $R_1(x)$ $R_3(z)$ $R_1(y)$ $w_1(x)$ $R_2(z)$ $w_2(y)$ $R_3(x)$ $w_3(z)$



Not Recoverable

$R_2(y)$ $R_1(x)$ $R_3(z)$ $R_1(y)$ $w_1(x)$ $R_2(z)$ $w_2(y)$ $R_3(x)$ $w_3(z)$



$T_1 T_2 T_3$
Conflict Serializable

Any Doubt ?

**THANK
YOU!**

